

Generative AI with OpenAI & ChatGPT

In 28
Minutes

Getting Started with Generative AI

In 28
Minutes

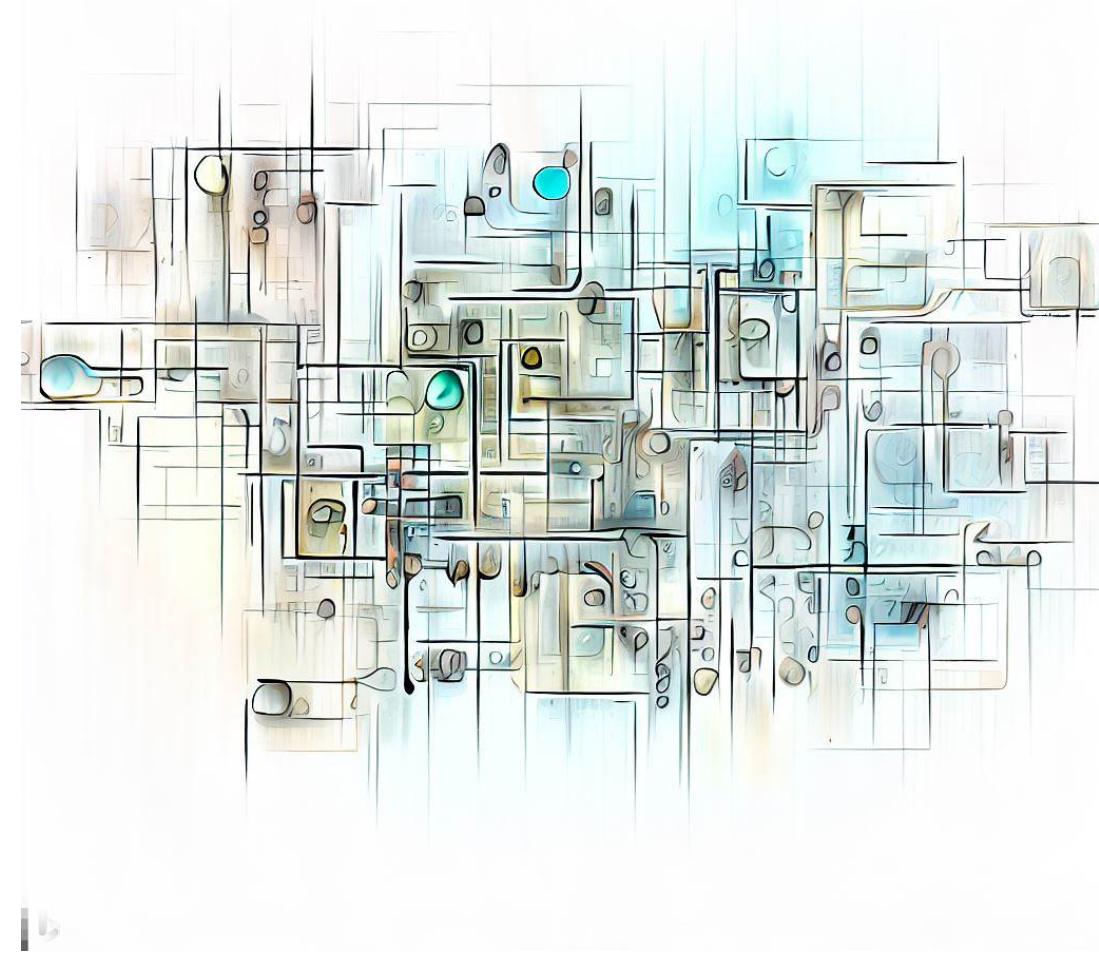
- Artificial Intelligence is in its **second golden phase** (IMHO):
 - **First Phase** - Early successes in the 1960s and 1970s
 - **Second Phase** - Last two decades or so driven by:
 - New algorithms
 - Massive datasets
 - Powerful computing hardware



Getting Started with Generative AI - 2

In 28
Minutes

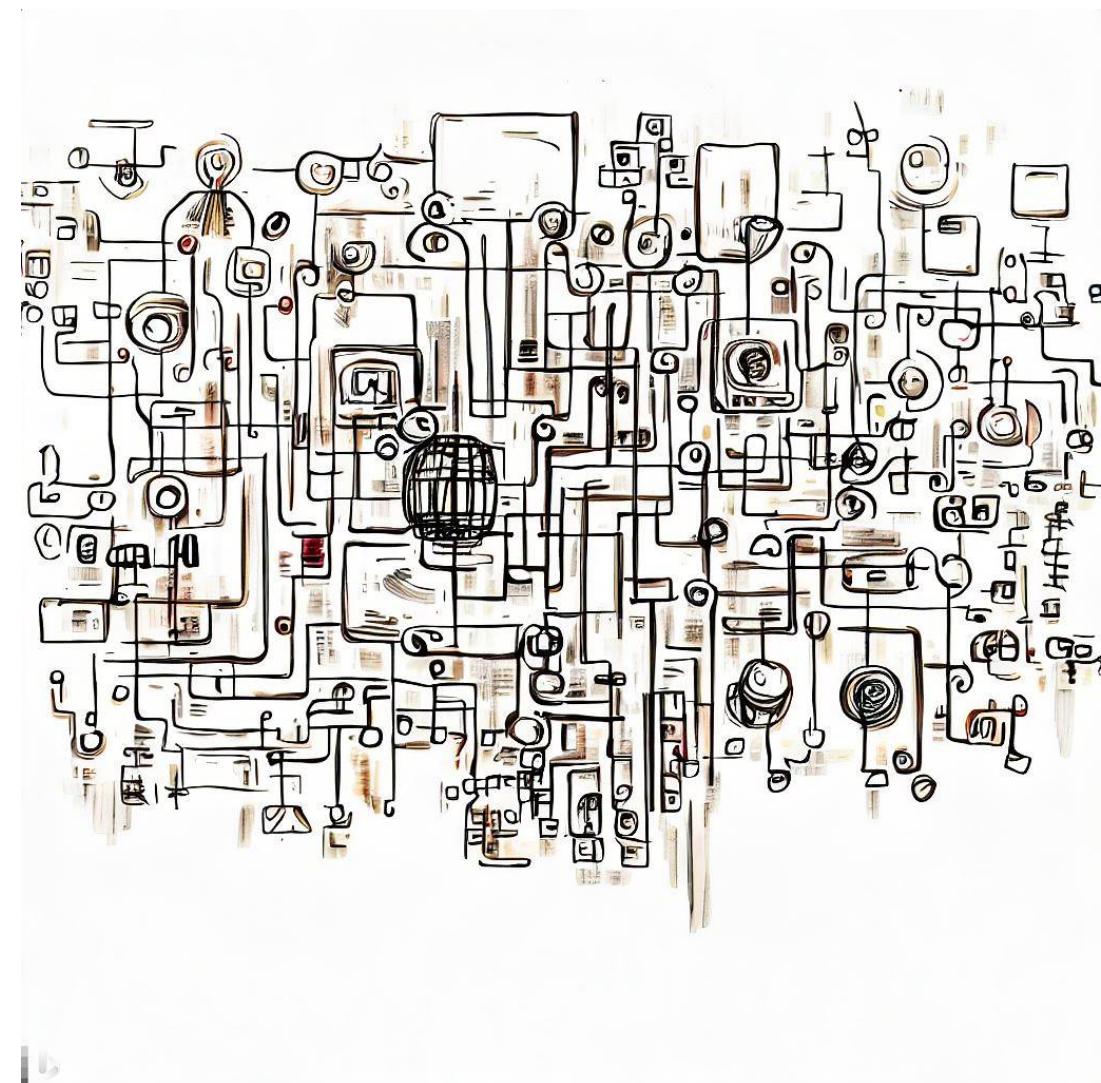
- HOWEVER, building AI solutions has NOT been easy:
 - Scarcity of skills
 - Need to create/manage massive datasets
 - Complex infrastructure needs
 - Complexity of managing AI models



Getting Started with Generative AI - 3

In 28
Minutes

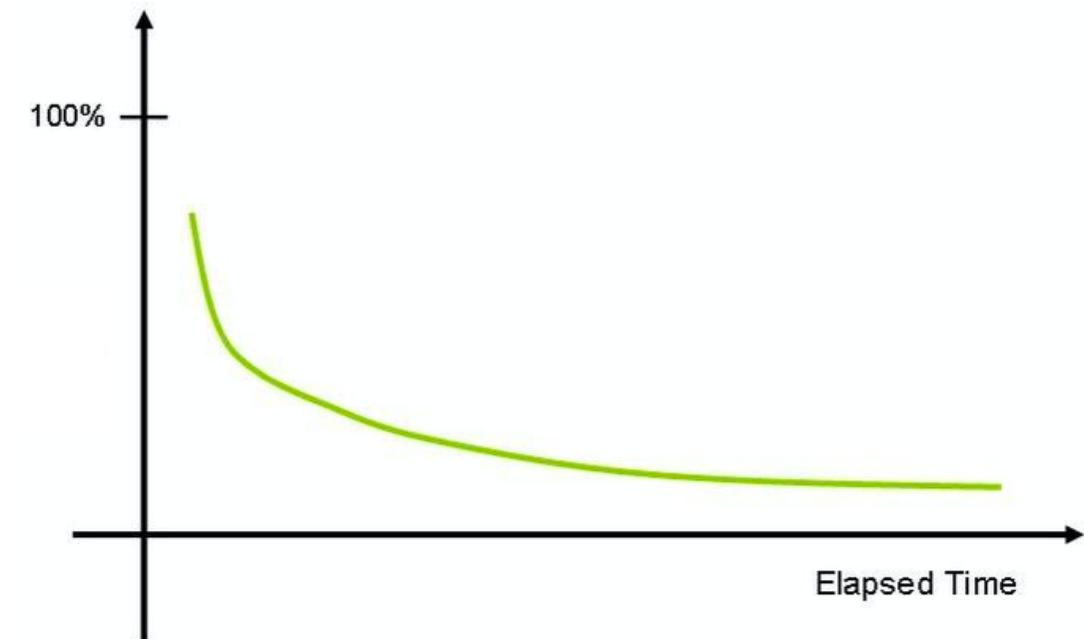
- How can you use AI WITHOUT AI skills/complex infrastructure?
 - Use low-code AI platforms (AutoML)
 - Use pre-trained models (especially Generative AI models)
 - OpenAI offers a plethora of solutions around Generative AI:
 - ChatGPT: Most famous Chatbot
 - OpenAI API: Easy to consume APIs
 - OpenAI Playground: Play with OpenAI APIs
- Our Goal : Help you understand AI, ML and Generative AI while exploring Generative AI solutions offered by OpenAI



How do you put your best foot forward?

In 28
Minutes

- Learning Gen. AI can be tricky:
 - Lots of new terminology
 - Lots of services
- As time passes, we forget things
- How do you **improve your chances of remembering things?**
 - Active learning - think and make notes
 - Review the presentation once in a while



Our Approach

In 28
Minutes

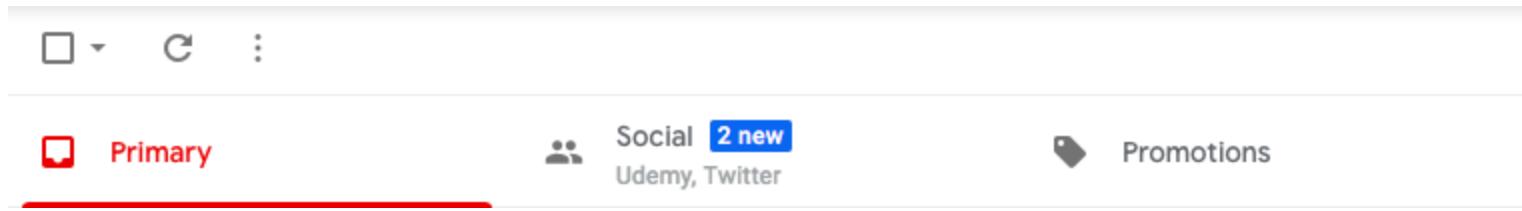
- Three-pronged approach to reinforce concepts:
 - Presentations (Video)
 - Demos (Video)
 - **Two kinds of quizzes:**
 - Text quizzes
 - Video quizzes
- (Recommended) Take your time. Do not hesitate to replay videos!
- (Recommended) Have Fun!



Artificial Intelligence

Artificial Intelligence - All around you

In 28
Minutes



- Self-driving cars
- Spam Filters
- Email Classification
- Fraud Detection

What is AI? (Oxford Dictionary)

In 28
Minutes

The theory and development of computer systems able to perform tasks normally requiring human intelligence, such as visual perception, speech recognition, decision-making, and translation between languages

What is AI? (in Simple Terms)

In 28
Minutes

- Goal of AI: Create machines that can simulate human-like intelligence and behavior
 - Play Chess
 - Play Go
 - Make purchase decisions
 - Drive a car
 - Write an essay!



Understanding Types of AI

In 28
Minutes

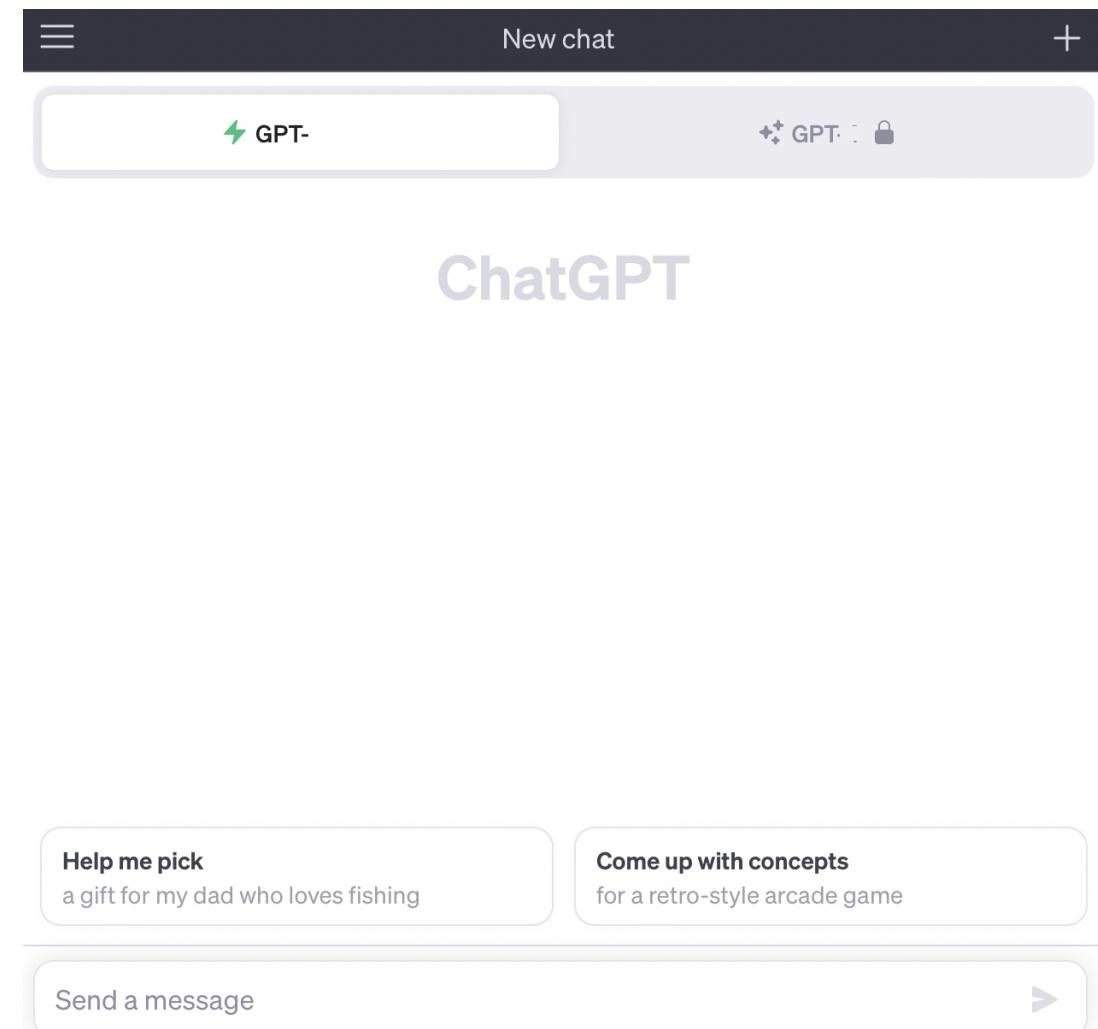
- **Strong artificial intelligence (or general AI):**
Intelligence of machine = Intelligence of human
 - A machine that can solve problems, learn, and plan for the future
 - An expert at everything
 - Including learning to play all sports and games!
 - Learns like a child, building on its own experiences
 - We are far away from achieving this!
 - Estimates: few decades to never
- **Narrow AI (or weak AI):** Focuses on specific task
 - Example: Self-driving car
 - Example: Playing Chess
 - Example: Predicting House Price



Playing with ChatGPT

In 28
Minutes

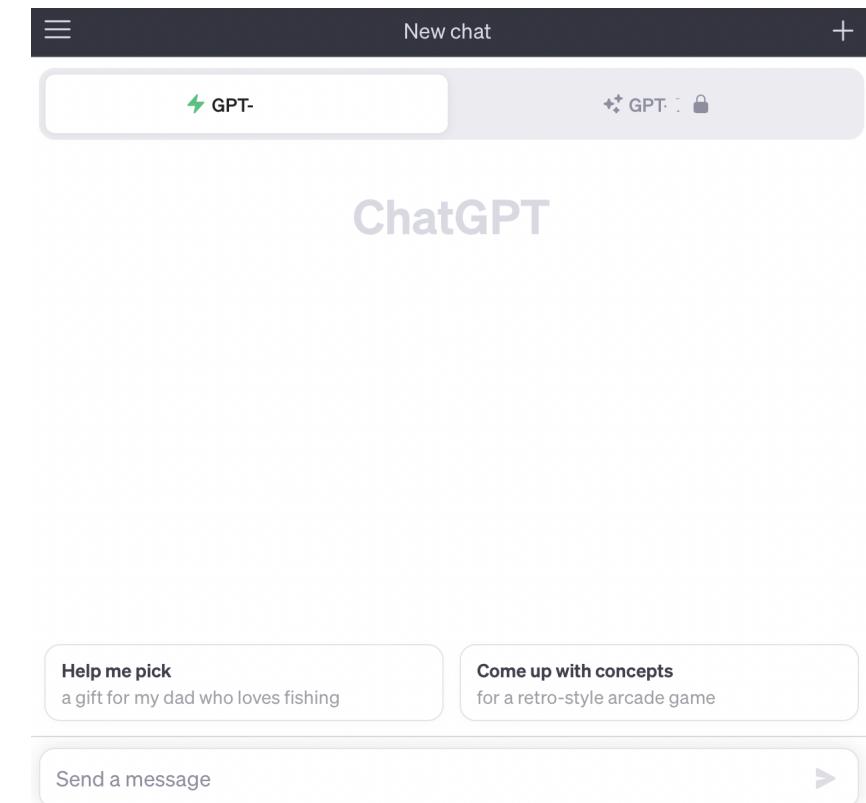
- ChatGPT: OpenAI's Generative AI Chatbot!
- A Demo of ChatGPT:
 - You are Lex Friedman. You are going to interview Sachin Tendulkar tomorrow. What are the FIRST FIVE questions that you are going to ask?
 - Act as Sachin Tendulkar. You are meeting Roger Federer. What questions would you ask?
 - Generate a bulleted list of items I need for an 15 day Everest Base Camp trek
 - I will be staying in tea houses on the trek. Can you update the list?



Playing with ChatGPT - Coding, Learning and Design

In 28
Minutes

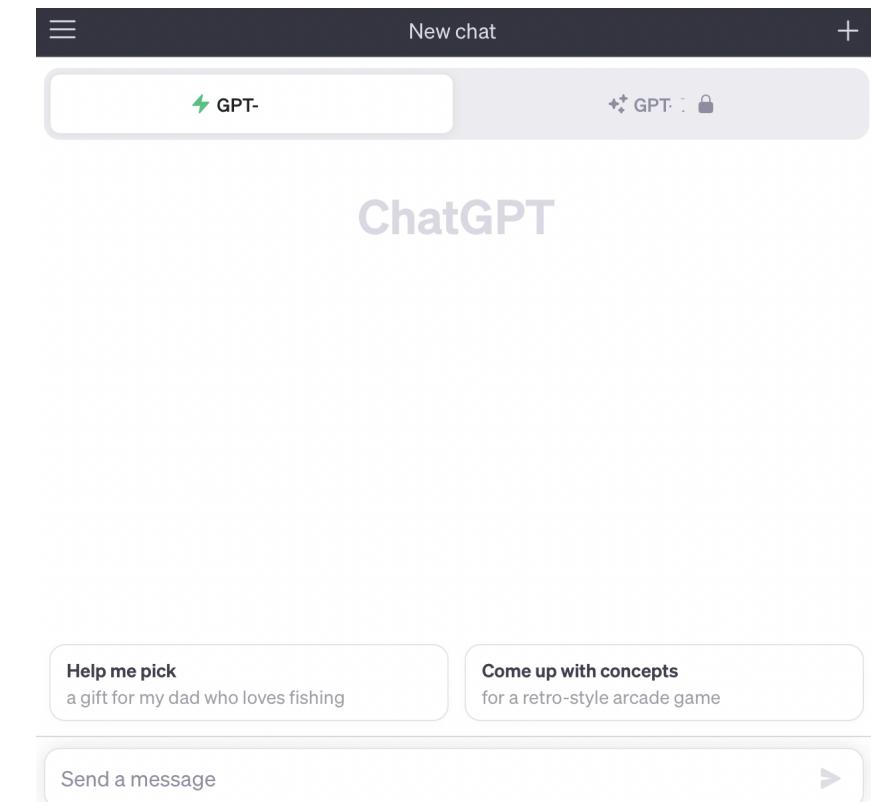
- Write a **Python function** to determine if a year is a leap year
- I'm learning Python for loop. Give me a **few exercises** to try?
- Can you **design a REST API** for todos? Give me an example request and response for each.
- I want to store information about courses, students, enrollments and reviews **in a relational table**. Can you suggest a structure?
- I like **learning concepts using a lot of examples**. What would be the books you would recommend to learn Design Patterns?



Playing with ChatGPT - Exploring Technology

In 28
Minutes

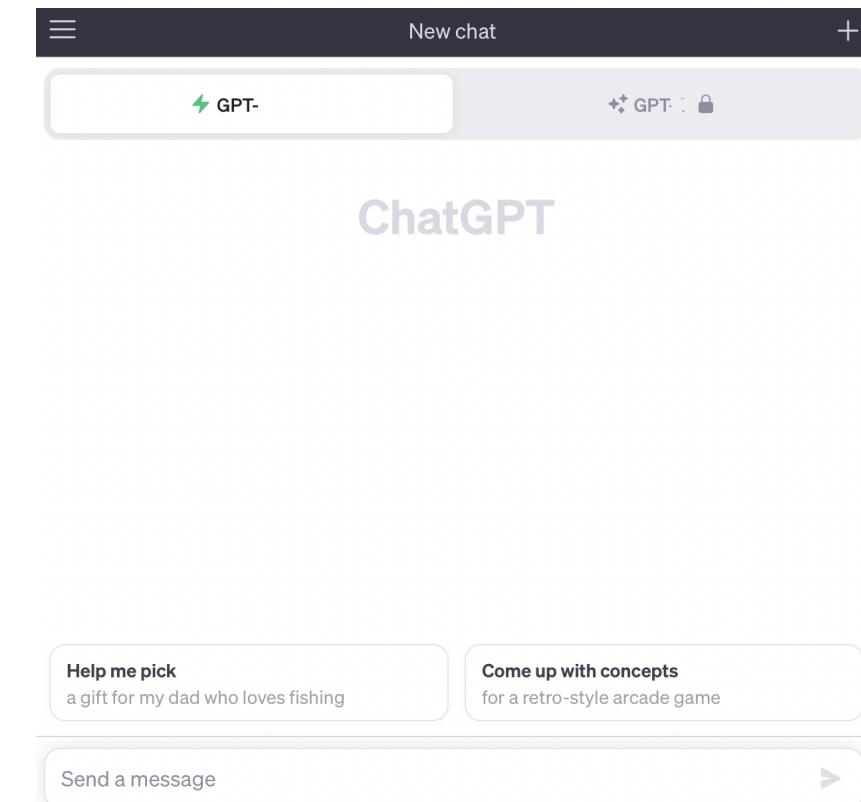
- Can you make list of **Top 10 technologies** that I might want to learn as a cloud engineer?
- For a new project, I'm considering React and Angular as front end frameworks.
 - Can you compare them and present the results in a **tabular format**? Feature/Factor in the column and the framework in the row.
- I like to **learn in a step-by-step approach** by breaking down complex concepts into smaller, more manageable parts.
 - How can I learn Docker? Give me a list of 10 step by step exercises I can begin with. Make sure you order them in increasing order of difficulty. Present the results in a table.



Playing with ChatGPT - Generating Ideas

In 28
Minutes

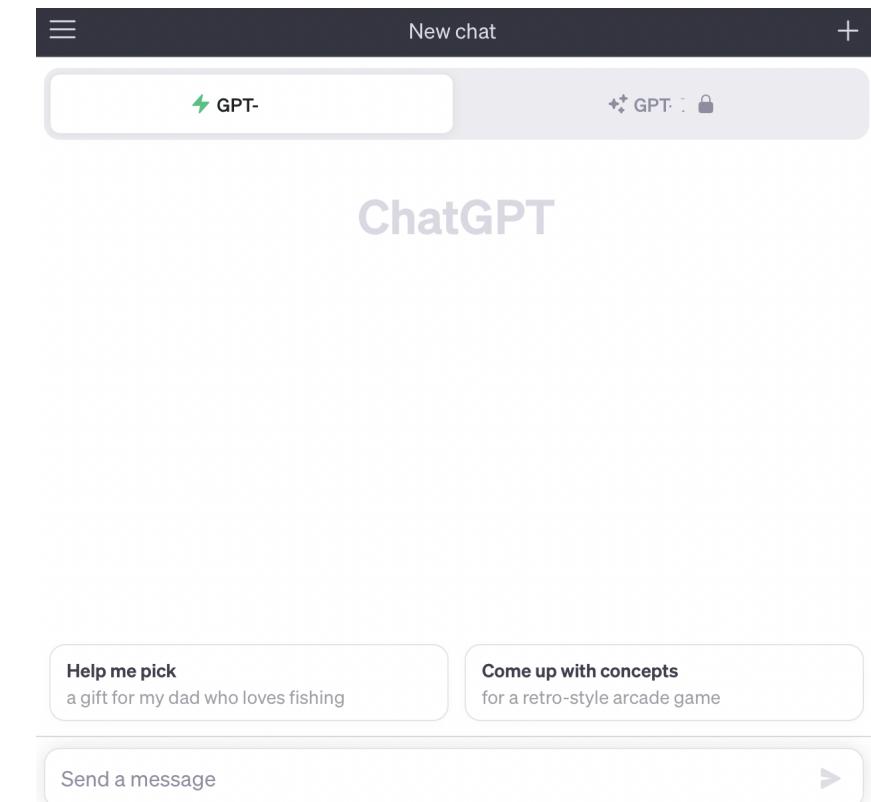
- I want to **Understand DevOps** and start exploring different DevOps Tools. Can you suggest a few prompts I can start with?
- How can I **make it easy for a beginner** to learn a new programming language / tool / technology / framework / methodology? One example is to **use analogies**. Can you think about other similar things I can do. Make sure you order them in decreasing order of importance. Rate importance from 1-100. 1 being least important. Present the results in a table.



Playing with ChatGPT - Observations

In 28
Minutes

- ChatGPT MAY display inaccurate or offensive information
- BUT it does provide a lot of value if you understand its limitations and know when/how to use it
- **How does ChatGPT work?**
 - Artificial Intelligence
 - Machine Learning
 - Generative AI
 - Large Language Models
 - Foundation Models



AI vs ML vs Generative AI

In 28
Minutes

- Goal of AI: Create machines that can simulate human-like intelligence and behavior
 - What is ML?
 - How does Generative AI fit in?
- Let's get started on a Journey!



Machine Learning vs Traditional Programming

- **Traditional Programming:** Based on Rules

- IF this DO that
- Example: Predict price of a home
 - Design an algorithm taking all factors into consideration:
 - Location, Home size, Age, Condition, Market, Economy etc

- **Machine Learning:** Learning from Examples (NOT Rules)

- Give millions of examples
- Create a Model
- Use the model to make predictions!

Home size (Square Yds)	Age	Condition (1-10)	Price \$\$\$
300	10	5	XYZ
200	15	9	ABC
250	1	10	DEF
150	2	34	GHI

0 0 0 0 0 0 0 0 0 0 0 0 0 0
1 1 1 1 1 1 1 1 1 1 1 1 1 1
2 2 2 2 2 2 2 2 2 2 2 2 2 2
3 3 3 3 3 3 3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6 6 6 6 6 6 6
7 7 7 7 7 7 7 7 7 7 7 7 7 7
8 8 8 8 8 8 8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9 9 9 9 9 9 9

Machine Learning Fundamentals - Scenarios

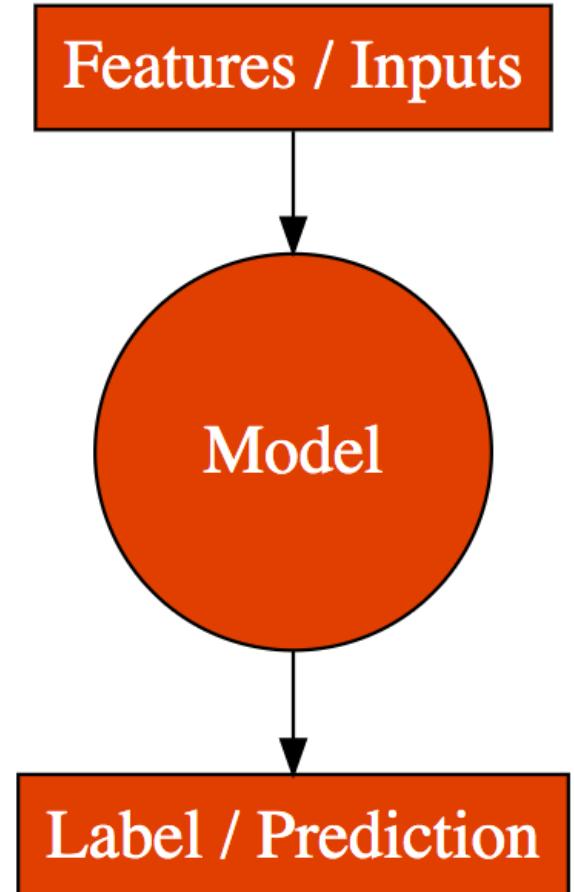
In 28
Minutes

Scenario	Solution
Categorize: Building a computer system as intelligent as a human. An expert at everything (all sports and games!)	Strong AI
Categorize: Building a computer system that focuses on specific task (Self-driving cars, virtual assistants, object detection from images)	Narrow AI (or weak AI)
Category of AI that focuses on learning from data (examples)	Machine learning
How is ML different from traditional programming?	Traditional Programming: Rules. Machine Learning: Examples

Machine Learning - Making Prediction

In 28
Minutes

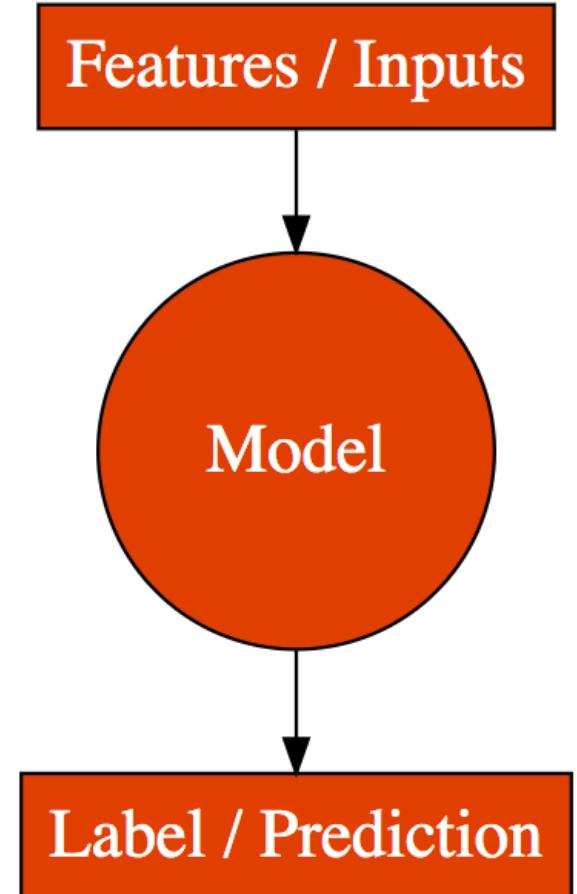
- **Goal:** Make a Good Prediction
 - Give inputs to a model
 - Model returns the prediction
 - Inputs are called **Features**
 - Prediction is called **Label**
 - **Example:** House Price Prediction Model
 - **Label:** price
 - **Features:**
 - area: Total area of house (m^2)
 - rooms: No. of rooms
 - bedrooms: No. of bedrooms
 - furniture: Is it furnished?
 - floor: Which floor?
 - age: How many years?
 - balcony: has balcony or not
 - garden: has garden or not



Machine Learning - Features and Labels - Examples

In 28
Minutes

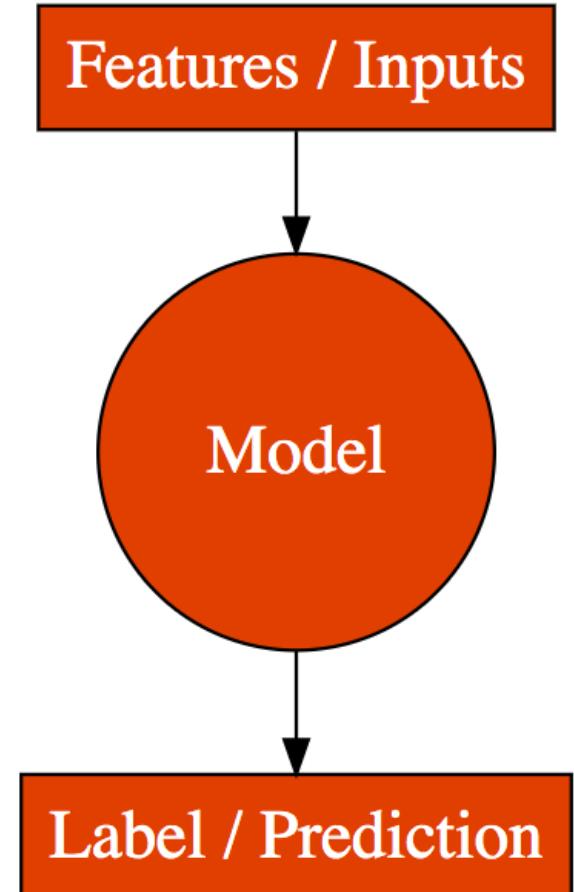
- Used Car Price Prediction Model
 - **Label:** price
 - **Features:** manufacturer, year, model, age, condition, cylinders, location
- Spam Email Classification Model
 - **Label:** isSpam
 - **Features:** sender, subject, content
- Grant a Loan Model
 - **Label:** shouldWeGrantALoan
 - **Features:** doesOwnCar, doesOwnRealEstate, creditScore, isMarried, doesHaveChildren, totalIncome, totalCredit



Machine Learning - Prediction Possibilities

In 28
Minutes

- **Numeric value:** Label is a numeric value with a range of possibilities => **Regression**
 - Used Car Price Prediction
 - House Price Calculation
 - Predicting sea level
 - How much will it rain tomorrow?
- **Limited Possibilities:** YES or NO, 0 or 1, Type 1 or Type 2 or Type 3 => **Classification**
 - Spam Email, Grant a Loan, Determine the type of cloud
 - Will it rain today?
- **Summary:**
 - **Classification:** Predicting category
 - **Regression:** Predicting numeric value



Machine Learning - Making Predictions - Scenarios

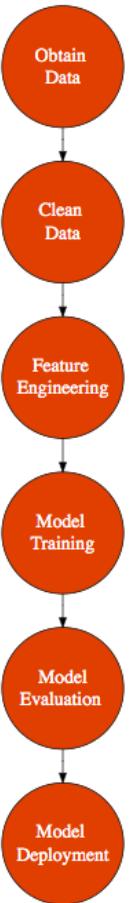
In 28
Minutes

Scenario	Solution
Categorize into features and labels for house price prediction: price, area, rooms, age	price is label. Others can be features
Categorize into features and label for used vehicle price prediction: manufacturer, year, model, age, condition, cylinders, location, price	price is label. Others can be features
Categorize: Used Car Price Prediction	Regression
Categorize: Spam Email Identification	Classification
Categorize: Predict amount of rainfall in the next year	Regression
Categorize: Should we grant a loan?	Classification
Categorize: Identify the type of vehicle in an image	Classification
Categorize: Find a specific dance form in a video	Classification

Creating Machine Learning Models - Steps

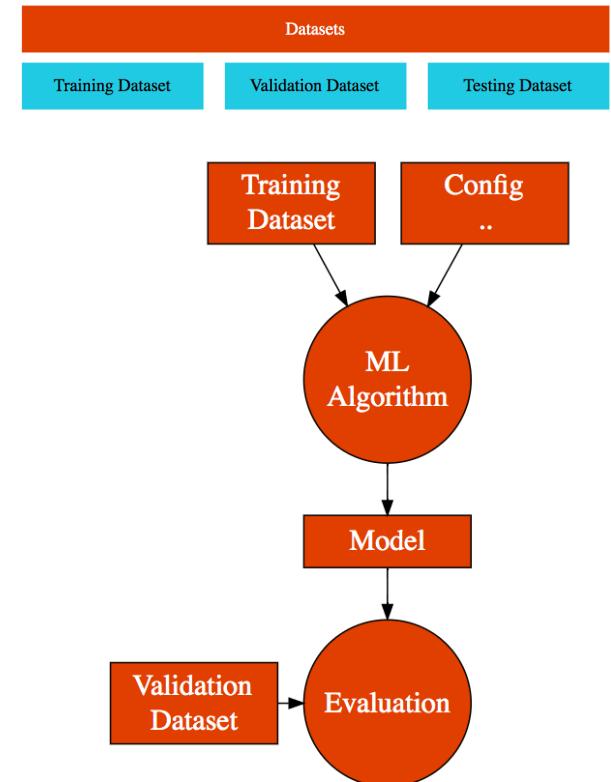
In 28
Minutes

- 1: Obtain Data
- 2: Clean Data
- 3: Feature Engineering: Identify Features and Label
- 4: Create a Model using the Dataset (and the ML algorithm)
- 5: Evaluate the accuracy of the model
- 6: Deploy the model for use



Understanding Machine Learning Terminology

- **Process**
 - **Training:** The process of creating a model
 - **Evaluation:** Is the model working?
 - **Inference:** Using model to do predictions in production
 - **Dataset:** Data used to create, validate & test the model
 - **Features:** Inputs
 - **Label:** Output/Prediction
 - **Dataset Types**
 - **Training Dataset:** Dataset used to create a model
 - **Validation Dataset:** Dataset used to validate the model (and choose the right algorithm) - Model Evaluation
 - **Testing Dataset:** Dataset used to do final testing before deployment



ML Stages and Terminology - Scenarios

In 28
Minutes

Scenario	Solution
Determine Stage: You remove data having null values from your dataset	Clean Data (Data Preparation)
Determine Stage: Normalize or split data into multiple features	Feature Engineering
Determine Stage: You evaluate the accuracy metrics of a model	Model Evaluation
Terminology: Using model to do predictions in production	Inference
Terminology: The process of creating a model	Training
Terminology: Dataset used to (train) or create a model	Training Dataset
Terminology: Dataset used to evaluate a model	Validation Dataset
Terminology: Dataset used to do final testing before deployment	Testing Dataset

THE AI Turmoil

In 28
Minutes

- **Quotes:**

- I am really quite close, I am very close, to the cutting edge in AI and it scares the hell out of me - **Elon Musk**
- The development of full artificial intelligence could spell the end of the human race. It would take off on its own, and re-design itself at an ever-increasing rate. Humans, who are limited by slow biological evolution, couldn't compete and would be superseded. - **Stephen Hawking**

- No one knows the truth:

- Most predictions about AI turned false in the last few decades!

- What's the pragmatic way to think?

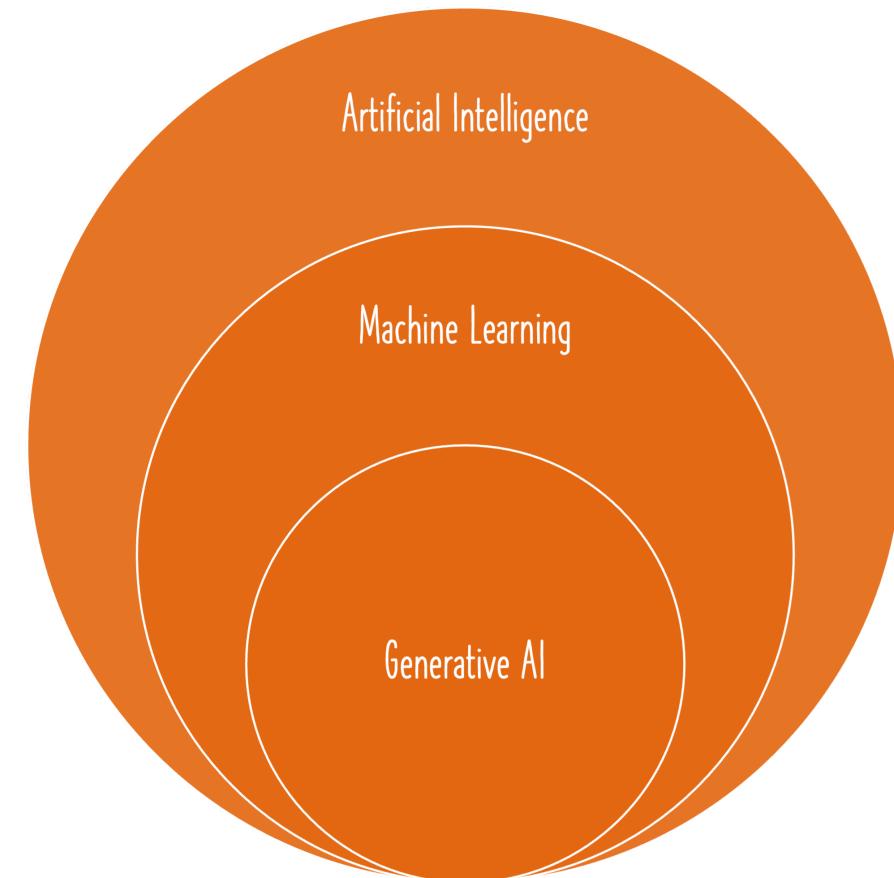
- Don't fear AI
- Learn to make the best use of it



Generative AI - How is it different?

In 28
Minutes

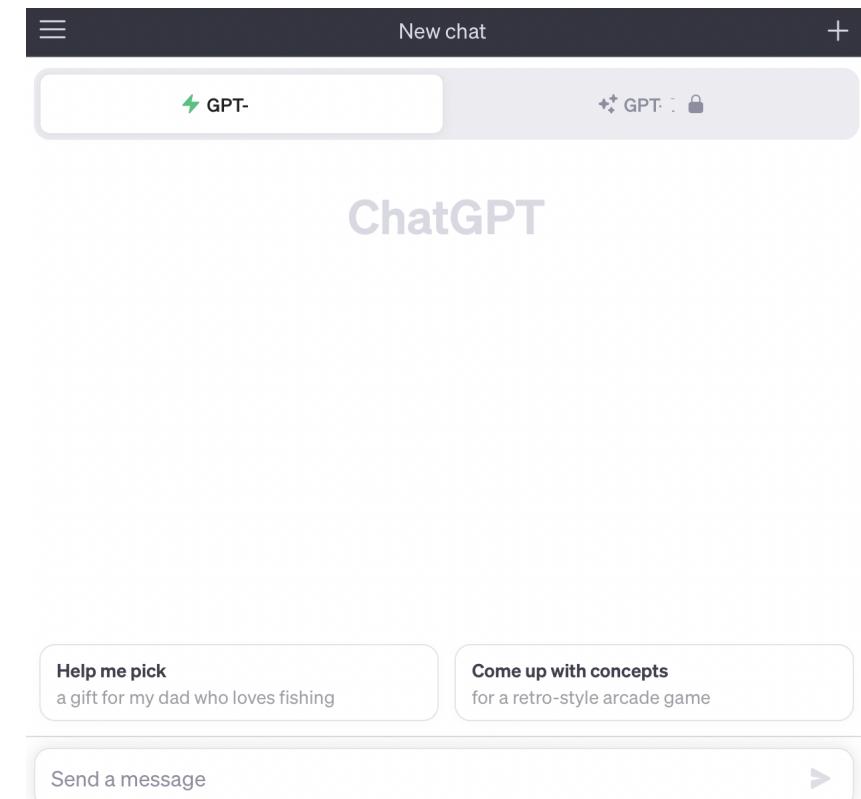
- **Artificial Intelligence:** Create machines that can simulate human-like intelligence and behavior
 - **Machine Learning:** Learning from examples
 - **Generative AI:** Learning from examples to create new content



Generative AI - Generating New Content

In 28
Minutes

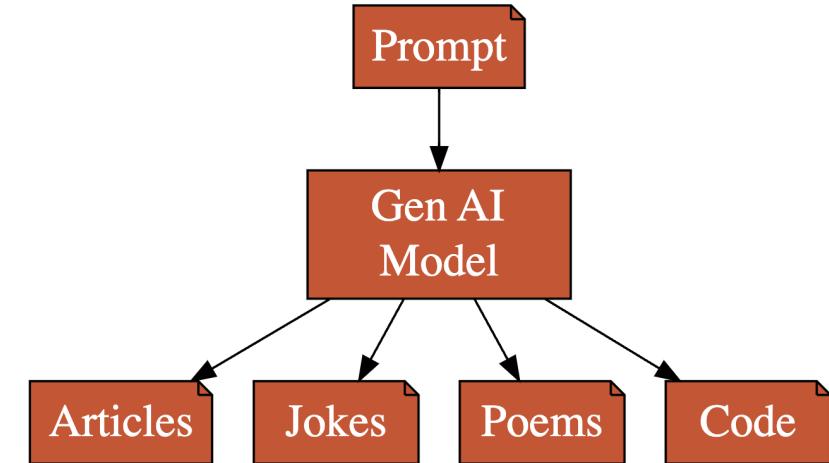
- **Goal:** Generating New Content
 - Instead of making predictions, Generative AI focuses on creating new data samples
 - **Examples:**
 - **Text Generation:** Writing e-mails, essays & poems. Generating ideas.
 - **Writing Code:** Write, debug & analyze programs
 - **Images Generation:** Creating paintings, drawings, or other forms of images
- How else is Generative AI different?
 - Let's find out!



Generative AI - Needs Huge Volumes of Data

In 28
Minutes

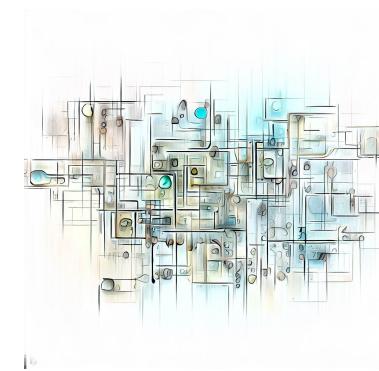
- **Generative AI models:** Statistical models that learn to generate new data by analyzing existing data
 - More data analyzed => Better new data similar to existing data
 - Example: GPT-3 model was trained on a dataset of 500 billion words of text
- **Datasets used include:**
 - Images, text and code scraped from the open web:
 - Wikipedia
 - Books
 - Open source code (syntax of programming languages and the semantics of code)
 - Conversations



Generative AI - Uses Self Supervised Learning

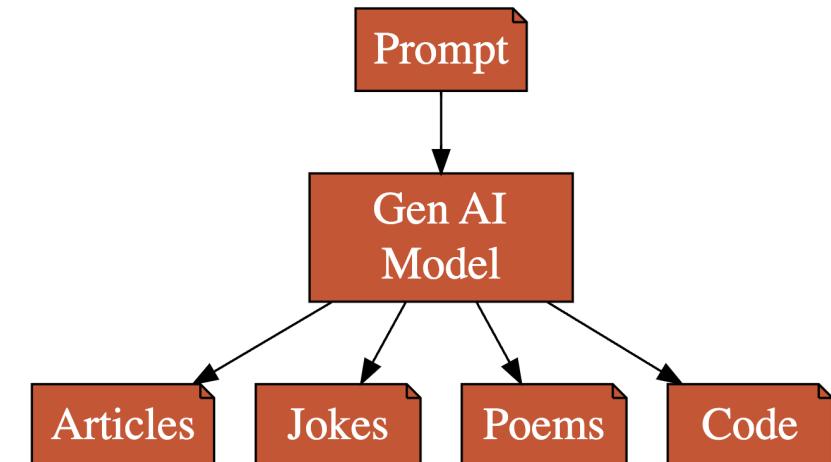
In 28
Minutes

- **Self-supervised learning:** Model learns from the data itself
 - WITHOUT requiring explicit labels or annotations
- How does this work?
 - Example for text model:
 - 1: Model tries to predict next word based on preceding words:
 - Model is given example sentence: "The sun is shining and the sky is __."
 - Model predicts the missing word
 - 2: Model's predicted word is compared to the actual word that comes next:
 - Learns from its mistakes and adjusts its internal representations
 - Neural Networks, Loss Calculation, Backpropagation etc..
 - 3: Repeated for all text from training dataset
 - Model captures the relationships between words, contextual cues, and semantic meanings:
 - If prompted with "The sun is shining and the sky is," the model might generate:
 - "The sun is shining and the sky is clear."
 - "The sun is shining and the sky is blue."
 - "The sun is shining and the sky is filled -- with fluffy clouds."



Key Step In Generative AI For Text - Next Word

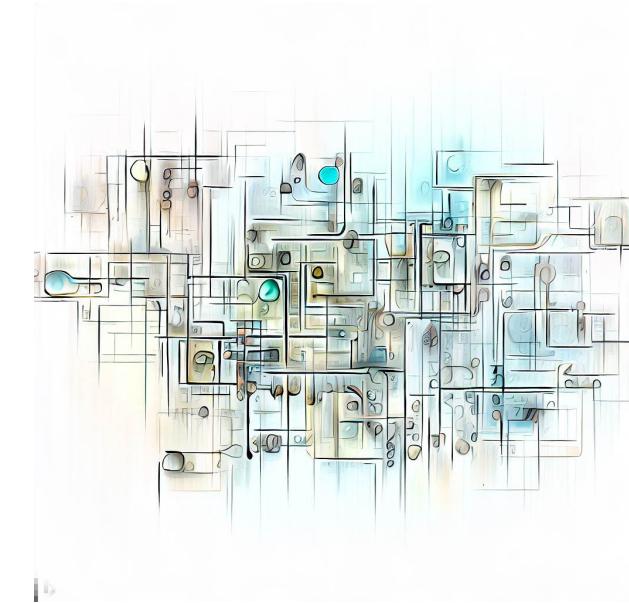
- A key step in Generative AI For Text is **predicting** the next word
- During training, text based Generative AI models **learn the probability** that a word might occur in a specific context
 - Context: "The cat sat on the"
 - **Example probabilities for next word:**
 - "mat": 0.4, "table": 0.2, "chair": 0.2, "moon": 0.1
 - Model **might** choose the highest probable word and go on to predict subsequent words
 - **HOWEVER**, you can **control** which of the words is chosen by controlling a few parameters!
 - temperature, top_k, top_p etc!



Generative AI Text - Uses Tokens instead of Words

In 28
Minutes

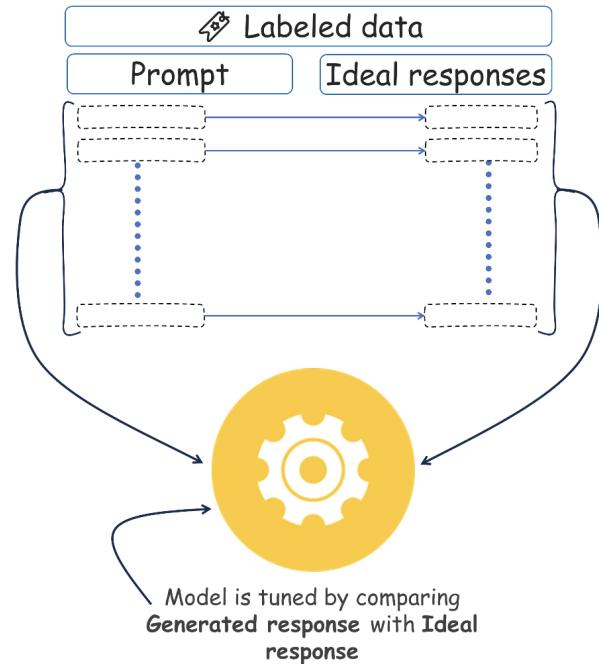
- **TOKEN:** A unit of text that might be a word
 - BUT it can be a sub word, punctuation mark, a number, ..
 - **Why Tokens?**
 - Tokens are **more consistent** than words
 - Words can have multiple meanings, depending on the context
 - "bank" might mean financial institution or a river bank
 - Tokens are more consistent
 - Example tokens: bank_river, bank_financial or light_verb, light_noun, ..
 - Tokens are **smaller** and more manageable
 - Tokens are **more efficient** to process
 - Because tokens are consistent, it's easy for models to learn relationships and things like parts of speech
- **Generative AI For Text Models:**
 - Understand relationships between ~~Words~~ **Tokens**
 - Good at predicting Next ~~Word~~ **Token!**
 - Have a **token limit** on context and generated text
 - Example: 1,024 tokens or 4,096 tokens



Generative AI Text - Uses SFT

In 28
Minutes

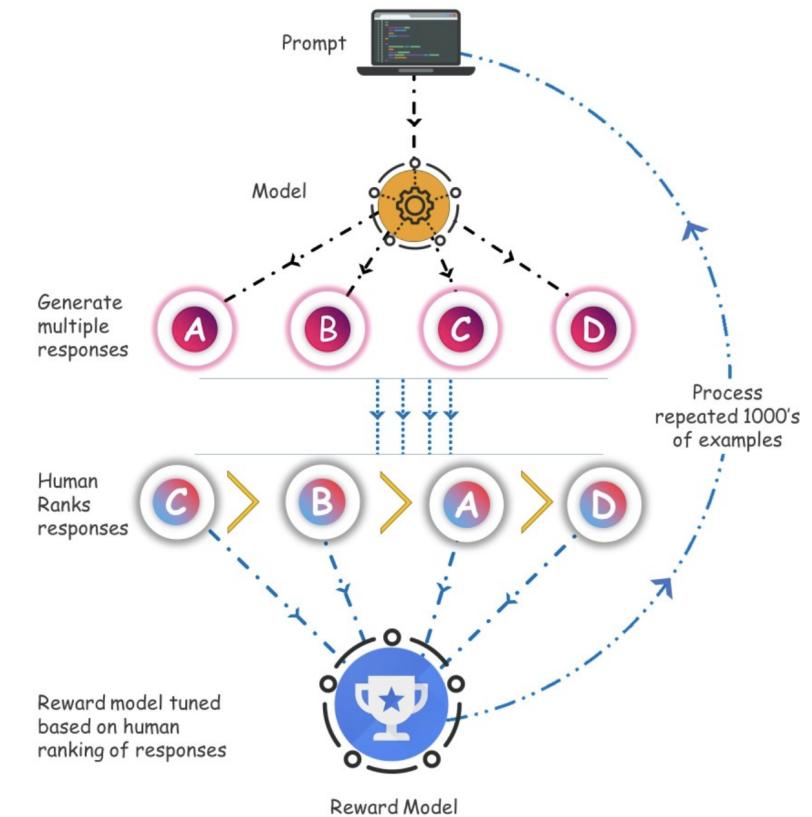
- After basic training, Gen. AI Model can predict **NEXT WORD** in a sequence based on contextual information
 - Given: "My favorite sport is"
 - Model picks a probable word (basketball:20%, soccer:18%, cricket:10%)
 - Given: "A question?"
 - Model might follow up with "Another Question?"
- **HOW to make model to respond to questions with answers?**
 - Given a question, how to make the model give an answer
 - **Solution:** Supervised Fine-Tuning
- Model is trained with **Labeled Data**
 - 1000s of Prompt and Ideal Response combinations
 - Model learns to respond to a question with an answer
 - Surprisingly less number of prompts add this capability



Generative AI Text - Uses RLHF

In 28
Minutes

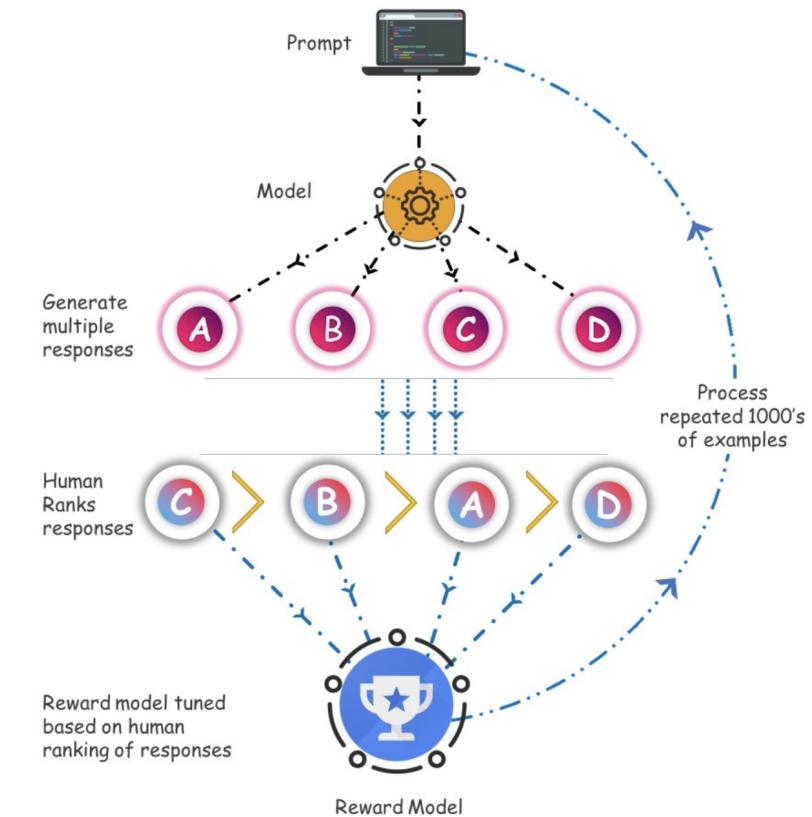
- How to make the model understand human values and preferences?
 - Models don't inherently understand human values, ethics, or preferences
 - Models can sometimes generate content that is inappropriate, biased, or conflicts with human values
 - How can we avoid this?
 - Solution: Reinforcement Learning from Human Feedback (RLHF)
- STEPS:
 - 1: Create a Reward Model that understands human values and preferences
 - 2: Tune Generative AI Model using Reward Model



Generative AI Text - Uses RLHF - Step 1

In 28
Minutes

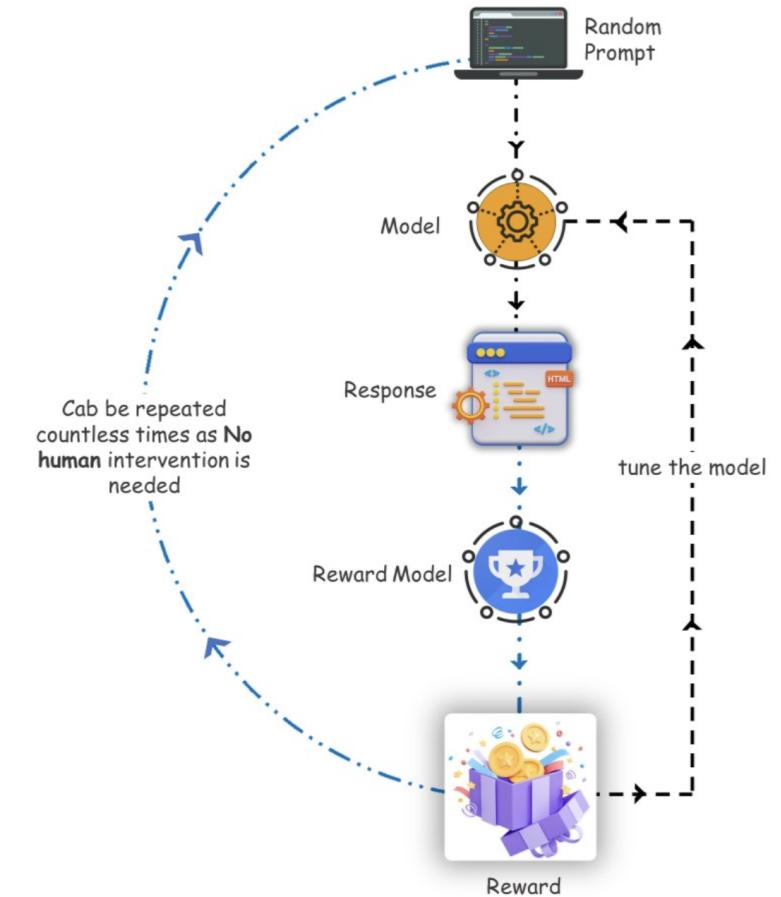
- **Goal:** Build a Reward Model that understands human values and preferences
- **How?**
 - 1: Use Generative AI model to generate multiple responses for a prompt
 - 2: Diverse set of human evaluators rank the responses
 - 3: Tune Reward Model based on the responses from human evaluators
 - Repeat 1, 2, 3 for thousands of prompts
- **Result:** Reward Model understands human values and preferences (**generates a Reward!**)
 - It will be used later to tune the responses from the Generative AI model



Generative AI Text - Uses RLHF - Step 2

In 28
Minutes

- **Goal:** Tune Generative AI Model using Reward Model
- **How?**
 - 1: Feed a prompt into Generative AI Model to generate a response
 - 2: Calculate Reward using Reward Model
 - 3: Tune Generative AI Model based on the evaluation
 - Repeat 1, 2, 3 for millions of prompts (automated)
- **Result:** Generative AI Model understands human values and preferences!
- **REMEMBER:** You don't need to understand everything about SFT and RHLF



Predictive Machine Learning vs Generative AI

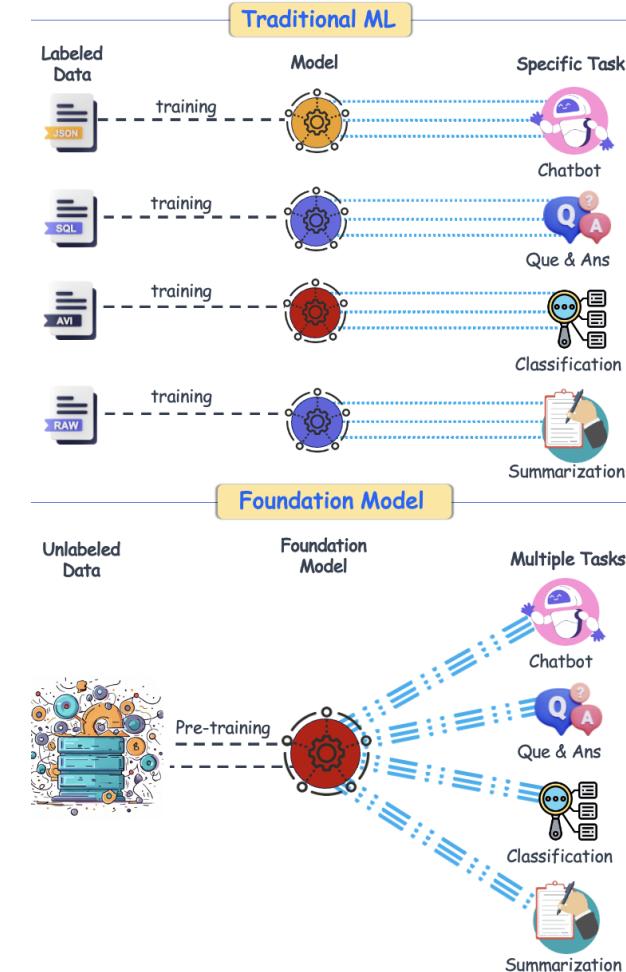
In 28
Minutes

Feature	Predictive Machine Learning	Generative AI
Goal	Make a Good Prediction	Generating New Content
Input	Features	Prompt
Output	Prediction (Label)	New Content
Use Cases	House Price Prediction, Fraud Detection, and more	Text Generation, Code Generation, Music Composition, and more
Volume of Training Data	Requires substantial labeled data	Requires significant amount of data
Time needed for Training	Training time can vary based on data size and complexity	Training time can be substantial for complex models

Generative AI - Foundation Models and LLMs

In 28
Minutes

- **Traditional ML Models:**
 - Needed task specific training
 - Multiple tasks => Multiple trainings => Multiple models
- **Foundation Models: Pre-Trained Multi Task Models**
 - Trained once (called pre-training)
 - Same model can be used for multiple tasks
 - Chatbot
 - Classification
 - Summarization
 - Some models are multi modal as well: Text, video, audio, image...
- **Large Language Models: Focused on text**
 - **REMEMBER:** Subset of Foundation Models
 - Models that are trained on a lot of text to generate more text!



Generative AI - Foundation Models

In 28
Minutes

- **From Bottom to Top:**

- **Foundation Models:**

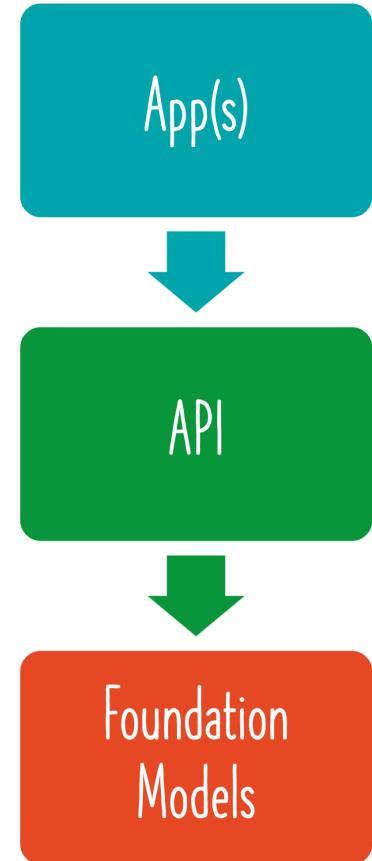
- OpenAI:
 - GPT (2, 3, 3.5, 4, 5, ..): Text, code and more..
 - DALL·E (1, 2, ...): Images
 - Open Source:
 - OpenLLaMA (Meta): Generate text, images, and code
 - Other Vendors: Google PaLM, ...

- **API:**

- OpenAI API
 - Azure OpenAI
 - Google Cloud PaLM API, ...

- **Applications**

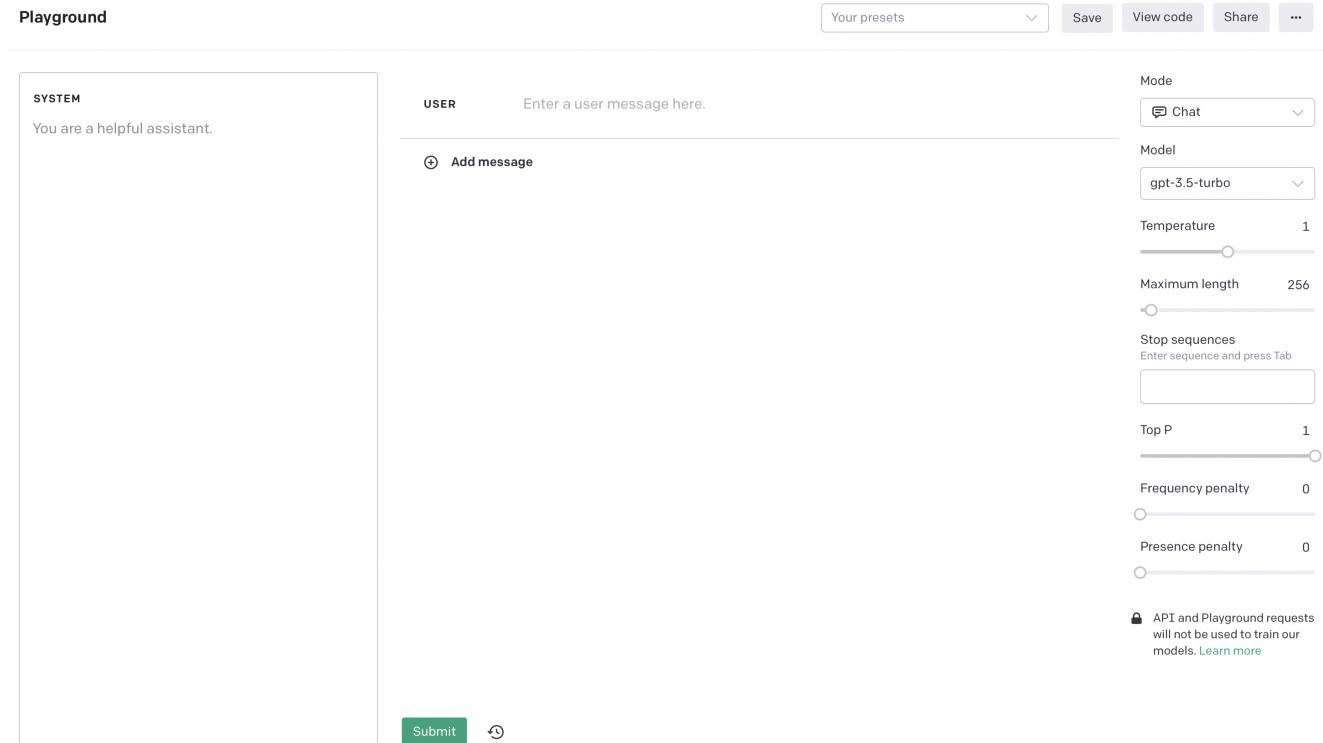
- ChatGPT
 - DALL·E
 - Bing Search, ...



Getting Started with OpenAI Playground

In 28
Minutes

- **OpenAI Playground:** Web based tool to experiment with the OpenAI API
 - <https://platform.openai.com/playground>
- **Great for experimenting:**
 - Play with different modes
 - Play with different models
 - Experiment with parameters
- Let's get started!



Exploring Text Features - Summarization

In 28
Minutes

- **Use Cases:**

- To summarize news articles from a website
- To summarize blog posts for a quick read
- To summarize technical documentation for a software product
- To summarize customer feedback for a business

- **Demos:**

- Summarize Reviews
- Create Hashtags
- Summarize Course Description
- Title generation for Course Description

The screenshot shows a user interface for text summarization. On the left, there's a text input field labeled "USER" containing the text: "What should be the right hashtags for this post: Why do I trek? What makes you walk for a few days continuously with just the most basic necessities on your back? I wish I knew the precise answer. The truth is - there are multiple layers to the answer. None satisfactory on its own but when everything is put together, it makes perfect sense (at least to me). Trekking has become a very important part of my life - new places, getting down to essentials, meeting great new people, ... Had a great time on the Kashmir Great Lakes trek. Do you trek? Why?" Below this is a "Hashtags:" section with an "Add message" button. On the right, there are several configuration options: "Mode" set to "Chat", "Model" (dropdown), "Temperature" (slider at 1), "Maximum length" (slider at 256), "Stop sequences" (text input), "Top P" (slider at 1), "Frequency penalty" (slider at 0), and "Presence penalty" (slider at 0).

Exploring Text Features - Classification

In 28
Minutes

• Use Cases:

- To classify customer feedback into different categories, such as positive, negative, or neutral.
- To classify news articles into different categories, such as politics, sports, or entertainment.
- To classify product reviews into different categories, such as helpful, unhelpful, or neutral.

• Demos:

- Sentiment (with JSON) identification
- Questions and Answers classification

USER
Text: This course is a fantastic introduction to the realm of microservices, docker, and kubernetes, ideal for newcomers seeking to explore these technologies. Alternatively, it serves as a valuable refresher on fundamental concepts for those already familiar with these subjects.
JSON:

ASSISTANT {
"sentiment": "Positive",
"explanation": "This review has a positive sentiment because the reviewer describes the course as a fantastic introduction to microservices, docker, and kubernetes. They also mention that it is ideal for newcomers and serves as a valuable refresher for those already familiar with the subjects."
}

+ Add message

Mode Chat

Model

Temperature 1

Maximum length 256

Stop sequences Enter sequence and press Tab

Top P 1

Frequency penalty 0

Presence penalty 0

Exploring Text Features - Extraction

In 28
Minutes

- **Use Cases:**

- To extract tech specs in a specific format
- Answer questions based on documentation

Playground

Your presets

Save View code Share ...

USER	Does the course cover AWS?
ASSISTANT	There is no mention of AWS in the description, so it is unclear whether the course covers AWS or not.
USER	Does the course cover Kubernetes?
ASSISTANT	Yes, the description explicitly states that the fourth part of the course covers the basics of Kubernetes and how to orchestrate microservices with Kubernetes.
USER	What is the title of the course?
ASSISTANT	The title of the course is not provided in the description.

SYSTEM
Answer questions based on this description.
Description:

Developing RESTful web services and REST API is fun. The combination of Java, Spring Boot, Spring Web MVC, Spring Web Services and JPA makes it even more fun. And its even more fun to create Microservices.

There are two parts to this course - RESTful web services and Microservices

Architectures are moving towards microservices. RESTful web services are the first step to developing great microservices. Java and Spring Boot, in combination with Spring Web MVC (also called Spring REST) makes it

Submit

Give us feedback

Exploring Text Features - Writing

In 28
Minutes

- **Use Cases:**

- Writing an Product announcement
- Writing an Ad copy
- Writing a Job post
- Writing an Email

The screenshot shows a user interface for generating text. On the left, under 'USER', the prompt is: 'Write an Article on DevOps' followed by 'I want to have the following sections:' and a list of topics: 'World Before Devops', 'Introduction to DevOps', 'Evolution of DevOps', 'Most Important DevOps Practices', 'Most Important DevOps Tools', and 'Where do I see DevOps Heading in Future?'. On the right, under 'ASSISTANT', the generated response is: 'World Before DevOps'. The text describes the traditional divide between development and operations teams before DevOps. At the bottom, there are 'Submit' and '↻' buttons, and a link to 'Give us feedback'. On the far right, there are various configuration options: Mode (set to Chat), Model, Temperature (set to 1), Maximum length (set to 278), Stop sequences (empty), Top P (set to 1), Frequency penalty (set to 0), Presence penalty (set to 0), and a note that API and Playground requests will not be used to train models.

Mode Chat

Model

Temperature 1

Maximum length 278

Stop sequences Enter sequence and press Tab

Top P 1

Frequency penalty 0

Presence penalty 0

API and Playground requests will not be used to train our models. [Learn more](#)

USER

Write an Article on DevOps

I want to have the following sections:

World Before Devops

Introduction to DevOps

Evolution of DevOps

Most Important DevOps Practices

Most Important DevOps Tools

Where do I see DevOps Heading in Future?

ASSISTANT

World Before DevOps

Before the concept of DevOps emerged, there was a clear divide between development teams and operations teams within organizations. Development teams were responsible for building and delivering software, while operations teams were responsible for deploying and maintaining that software in production environments. This traditional approach often led to inefficiencies, slow delivery cycles, and a lack of collaboration between the two teams.

Submit

↻

Give us feedback

Exploring Text Features - Ideation

In 28
Minutes

- **Use cases:**

- Generating new ideas
- Creative naming
- Get Advice
- Generate Interview questions

- **A few examples:**

- Generate a list of Top 10 DevOps and Cloud Trends
- Generate a name for an e-learning company focusing on teaching cloud
- List 5 best practices with respect to managing costs in the cloud
- List 10 interview question to ask a beginner to DevOps

The screenshot shows a user interface for generating text. On the left, a user input field contains "Generate a list of Top 10 DevOps and Cloud Trends". To the right, an "ASSISTANT" section displays a numbered list of seven items related to DevOps and Cloud Trends. The interface includes various configuration sliders and dropdowns on the right side.

USER	ASSISTANT
Generate a list of Top 10 DevOps and Cloud Trends	1. Kubernetes adoption: The use of Kubernetes for container orchestration is rapidly increasing as organizations embrace cloud-native architectures. 2. Serverless computing: The popularity of serverless computing is growing, enabling developers to focus on code without worrying about underlying infrastructure. 3. DevSecOps: Integrating security into the DevOps process is becoming critical as organizations aim to build secure applications and infrastructure. 4. Hybrid cloud solutions: Organizations are adopting hybrid cloud models, combining public and private clouds, allowing them to leverage the benefits of both. 5. AIOps: Artificial intelligence for IT operations is being utilized to automate and improve various aspects of infrastructure management and monitoring. 6. GitOps: The practice of using Git version control system for managing infrastructure and provisioning resources is gaining traction. 7. Infrastructure as Code (IaC): Organizations are increasingly adopting IaC tools like Terraform and CloudFormation to automate the provisioning and management of infrastructure.

Mode: Chat

Model:

Temperature: 1

Maximum length: 256

Stop sequences: Enter sequence and press Tab

Top P: 1

Frequency penalty: 0

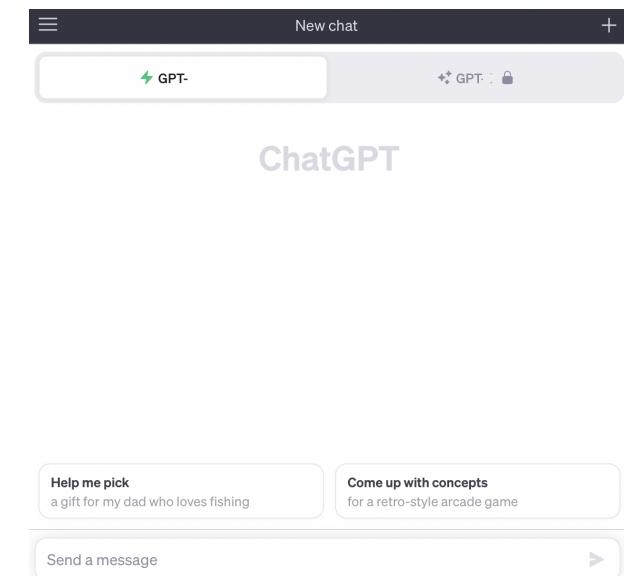
Presence penalty: 0

API and Playground requests will not be used to train our models. [Learn more](#)

Getting Started with Prompt Design

In 28
Minutes

- **Prompt:** A set of initial instructions provided to a foundation model as input
- **Prompt Design:** The process of crafting effective and precise prompts to achieve desired outputs from the language model
- **Why is Prompt Design Important?**
 - **Optimizing Model Performance:** A well-designed prompt can significantly impact the quality and relevance of the model's responses
 - **Getting the right response:** Leverage the full potential of foundation models, ensuring reliable, accurate, and contextually appropriate responses



Prompt Design - An Overview

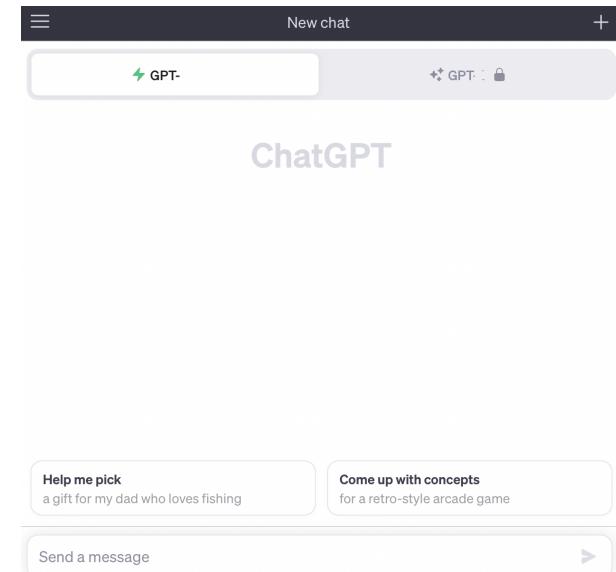
In 28
Minutes

- **Best Practices**

- **Clear instructions:** Avoid ambiguity or vagueness
 - Explain Docker in 100 words. Write the explanation so that a non technical person can understand.
- **Give examples (ZERO SHOT vs ONE SHOT vs FEW SHOT)**
- **Experiment to find the right prompt**
- **Consider using a framework (RTF, CTF, RASCEF, ...)**

- **Example:**

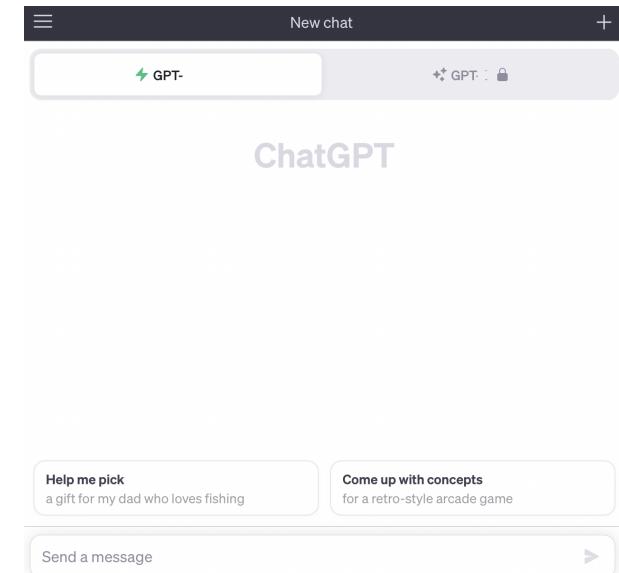
- **Context:** You are an DevOps and Cloud expert with 30 years of software experience. You are interviewing an expert for the position of an Cloud and DevOps architect.
- **Task:** What interview questions would you ask?
- **Example:**
 - I. Introduce yourselves



Prompt Design - Things to Try - An Exercise

In 28
Minutes

- Be **really specific** about the output you want
 - List 3 Advantages
 - Use JSON format
 - Specify the desired length of the output
- Define what to do if the model does not know the answer
 - If you are not confident, respond with "I don't know"
- Think step-by-step:
 - Ask for a chain of reasoning
- Ask the model to answer as if it was an expert:
 - "correct, high-quality, and written by an expert"



ZERO SHOT vs ONE SHOT - Example

In 28
Minutes

- **EXAMPLE 1: ZERO SHOT**

- Please choose the right answer:
- Question: Which of these is a programming language?
 - A) Docker
 - B) Python

- **EXAMPLE 2: ONE SHOT**

- Please choose the right answer:
 - Question: Which of these is a container orchestration tool?
 - A) Docker
 - B) Kubernetes
 - Answer: A is correct
 - Question: Which of these is a programming language?
 - A) Docker
 - B) Python
 - Answer:

▼ Prompt

Please choose the right answer:
Question: Which of these is a container orchestration tool?
A) Docker
B) Kubernetes
Answer: A is correct
Question: Which of these is a programming language?
A) Docker
B) Python
Answer:

▼ Response

— Markdown

B is correct

ZERO SHOT vs ONE SHOT vs FEW SHOT - Example

In 28
Minutes

- **EXAMPLES:**

- **EXAMPLE 1: ZERO SHOT**

- For the given order, return a JSON object
 - Order: A pizza and a pepsi

- **EXAMPLE 2: ONE SHOT**

- For the given order, return a JSON object
 - Order: A pizza and a pepsi
 - Output: {"pizza": 1, "pepsi": 1}
 - Order: A burger and a soda
 - Output:

- **Example 3: FEW SHOT**

- For the given order, return a JSON object
 - Order: A pizza and a pepsi
 - Output: {"pizza": 1, "pepsi": 1}
 - Order: A burger and 2 sodas
 - Output: {"burger": 1, "soda": 2}
 - Order: A burger, A pizza and 2 sodas
 - Output:

▼ **Prompt**

For the given order, return a JSON object
Order: A pizza and a pepsi
Output: {"pizza": 1, "pepsi": 1}
Order: A burger and 2 sodas
Output: {"burger": 1, "soda": 2}
Order: A burger, A pizza and 2 sodas
Output:

▼ **Response** Markdown

{"burger": 1, "pizza": 1, "soda": 2}

Simple Prompt Framework - RTF

In 28
Minutes

- **Role:** Define a persona for the user
- **Task:** Clearly state the specific task or question
- **Format:** Provide instructions on the desired format for the response
 - Example: a summary, pros and cons, a debate, or a step-by-step explanation.
- **Example:**
 - **Basic Prompt:** "Tell me about climate change."
 - **Improved RTF Prompt:**
 - **Role:** You're a climate scientist.
 - **Task:** Explain the causes and impacts of climate change.
 - **Format:** Provide a concise summary of key factors driving climate change and its consequences.



Role: You're a climate scientist.

Task: Explain the causes and impacts of climate change.

Format: Provide a concise summary of key factors driving climate change and its consequences.



Climate change is primarily driven by human activities, with the burning of fossil fuels like coal, oil, and natural gas being the chief contributor. This releases greenhouse gases, such as carbon dioxide and methane, into the atmosphere, trapping heat and leading to a rise in global temperatures. Deforestation and industrial processes also release additional greenhouse gases.

The impacts of climate change are far-reaching. Rising temperatures lead to more frequent and severe heatwaves, impacting human health and agriculture. Melting ice caps and glaciers contribute to rising sea levels, threatening coastal communities with increased flooding and erosion. Changes in precipitation patterns result in droughts, affecting water availability and agriculture. Additionally, shifts in ecosystems disrupt biodiversity and can lead to the extinction of vulnerable species.

Climate change exacerbates existing socio-economic challenges, particularly for marginalized communities who are more vulnerable to its effects. Adapting to these changes and mitigating further impacts requires collective global efforts to reduce emissions and transition towards sustainable, low-carbon alternatives.

Simple Prompt Framework - CTF

In 28
Minutes

- **Context:** Set the background or context
- **Task:** Clearly outline the specific task or question
- **Format:** Provide instructions on the desired format for the response
 - Example: a summary, pros and cons, a debate, or a step-by-step explanation.
- **Example:**
 - **Basic Prompt:** "Write about electric cars."
 - **Improved CTF Prompt:**
 - **Context:** In a world transitioning to sustainable energy.
 - **Task:** Explain the benefits and challenges of electric cars.
 - **Format:** Provide a balanced analysis, discussing environmental impact, technology, and adoption barriers.



Context: In a world transitioning to sustainable energy.

Task: Explain the benefits and challenges of electric cars.

Format: Provide a balanced analysis, discussing environmental impact, technology, and adoption barriers.



In a world transitioning to sustainable energy, electric cars present a promising solution to mitigate environmental impact. They produce zero tailpipe emissions, reducing air pollution and dependence on fossil fuels. This shift contributes significantly to combating climate change and improving urban air quality. Additionally, electric cars have lower operating costs due to the lower price of electricity compared to gasoline, and they require less maintenance as they have fewer moving parts.

However, challenges persist in the widespread adoption of electric cars. One major concern is the charging infrastructure. While it's expanding rapidly, especially in urban areas, it still lags behind the convenience of traditional gas stations, particularly in less populated regions. Additionally, the cost of electric vehicles (EVs) can be a barrier for many consumers, although advancements in battery technology are steadily reducing this gap. Furthermore, the environmental benefits are contingent on the energy mix used to generate electricity; if it's predominantly from fossil fuels, the emissions reduction from EVs may be less pronounced.

Prompt Framework - RASCEF

In 28
Minutes

- **Role:** Define the persona
- **Action:** Clearly state the action or the problem
- **Steps:** Outline the sequential or logical steps
- **Context:** Provide relevant background info
- **Examples:** Offer illustrative examples, if possible
- **Format:** Provide instructions on the desired format for the response
 - Example: a summary, pros and cons, a debate, or a step-by-step explanation.
- **Task:** Summarize the overall task or question, combining the elements above.



Role: You're an environmental researcher.
Action: Compare various renewable energy sources.
Steps: Analyze solar, wind, hydro, & geothermal energy.
Context: In a study on sustainable energy options.
Examples: Mention efficiency variations and geographical suitability.
Format: Prepare a chart with key attributes.
Task: Develop a comparative chart showcasing strengths and limitations of different renewable energy sources.



Certainly! Below is a comparative chart showcasing key attributes of different renewable energy sources: solar, wind, hydro, and geothermal energy.

Attribute	Solar Energy	Wind Energy	Hydro Energy	Geothermal Energy
Energy Efficiency	Moderate to High	High	High	High
Geographic Suitability	Suitable in Most Regions	Best in Coastal and Open Areas	Requires Suitable Topography and Flow	Suitable in Areas with Geothermal Activity
Resource Availability	Abundant	Abundant	Dependent on Water	Limited to Regions with Geothermal Activity
			Resources	Geothermal Activity
Environmental Impact	Low Greenhouse Gas Emissions	Low Impact on Wildlife	Alters Ecosystem and Fish Habitats	Potential for Land Subsidence
Land Use	Requires Space	Requires Open Space	Requires Large Area	Requires Minimal Land Footprint

Prompt Framework - RASCEF Example

In 28
Minutes

- **Basic Prompt:**

- "Compare different types of renewable energy."

- **Improved RASCEF Prompt:**

- **Role:** You're an environmental researcher.
- **Action:** Compare various renewable energy sources.
- **Steps:** Analyze solar, wind, hydro, & geothermal energy.
- **Context:** In a study on sustainable energy options.
- **Examples:** Mention efficiency variations and geographical suitability.
- **Format:** Prepare a chart with key attributes.
- **Task:** Develop a comparative chart showcasing strengths and limitations of different renewable energy sources.



Role: You're an environmental researcher.

Action: Compare various renewable energy sources.

Steps: Analyze solar, wind, hydro, & geothermal energy.

Context: In a study on sustainable energy options.

Examples: Mention efficiency variations and geographical suitability.

Format: Prepare a chart with key attributes.

Task: Develop a comparative chart showcasing strengths and limitations of different renewable energy sources.



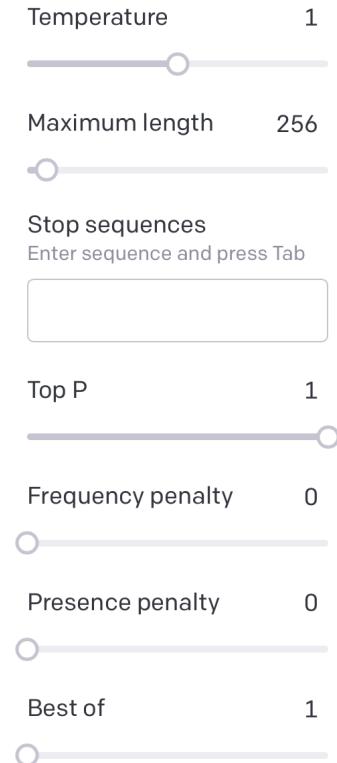
Certainly! Below is a comparative chart showcasing key attributes of different renewable energy sources: solar, wind, hydro, and geothermal energy.

Attribute	Solar Energy	Wind Energy	Hydro Energy	Geothermal Energy
Energy Efficiency	Moderate to High	High	High	High
Geographic Suitability	Suitable in Most Regions	Best in Coastal and Regions	Requires Suitable Topography and Flow	Suitable in Areas with Geothermal Activity
Resource Availability	Abundant	Abundant	Dependent on Water Resources	Limited to Regions with Geothermal Activity
Environmental Impact	Low Greenhouse Gas Emissions	Low Impact on Wildlife	Alters Ecosystem and Fish Habitats	Potential for Land Subsidence
Land Use	Requires Space	Requires Open Space	Requires Large Area	Requires Minimal Land Footprint

Getting Started with Parameters

In 28
Minutes

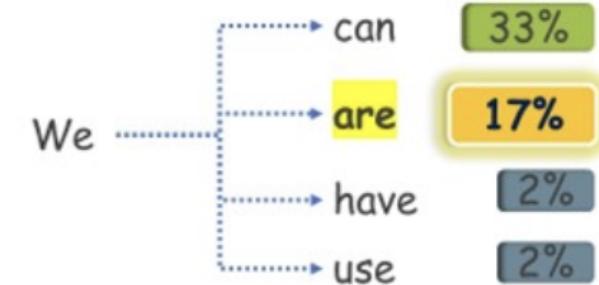
- How to customize the response from a model?
 - Answer: Configure Parameters
- Options:
 - Temperature
 - Maximum length
 - Top P
 - Frequency penalty
 - Presence penalty



Exploring Parameters

In 28
Minutes

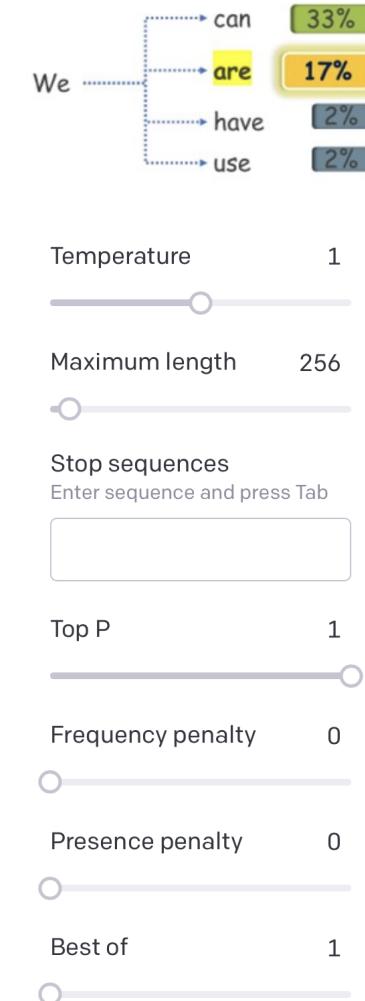
- Maximum length (`max_tokens`): How long do you want the response to be?
 - Specify maximum number of tokens in response
 - Remember: A token is approx. 4 characters. 100 tokens approximate to 60-80 words.
- Next set of parameters help you determine which token is chosen!
 - Example: A:0.4, B:0.2, C:0.1, D:0.05, E:0.02, F:0.01
 - Which of these should be chosen?
 - Temperature, Top P, Frequency Penalty, Presence Penalty
 - This is an Art (NOT a science)



Exploring Parameters - 2

In 28
Minutes

- A:0.4, B:0.2, C:0.1, D:0.05.. which of these should be chosen?
 - **Presence Penalty:** Penalize tokens if they are already present
 - Ex: token prob. = 0.4, pp = 0.1, previous occurrences = 2 => computed prob. - $0.4 - 0.1 = 0.3$
 - **Frequency Penalty:** Penalize tokens based on frequency in text so far
 - Ex: token prob. = 0.4, fp = 0.1, previous occurrences = 2 => computed prob. - $0.4 - 2 * 0.1 = 0.2$
 - **Temperature:** How random should be the output?
 - Higher values => more randomness and more creativity
 - Lower values => lesser randomness
 - **Top P:** What is the (cumulative) probability limit ?
 - Define the cumulative probability cutoff for selecting tokens
 - Lower value => less random responses. Higher value => more random responses.
 - Example: top_p value is 0.6 => Next token is either A or B
 - OpenAI recommends altering temperature or top_p but not both.
 - Example Scenarios:
 - Find Capital City of India: use low values
 - Write a creative essay: use high values



Getting Started with OpenAI API

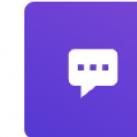
In 28
Minutes

- **OpenAI API:** Integrate Generative AI into your apps

- Understand and generate natural language and code
- Generate and edit images
- Convert speech into text
- Fine tune models

- **Models:**

- gpt-4, gpt-3.5-turbo, text-davinci-003, ..



GPT

Learn how to generate text



Embeddings

Learn how to search, classify, and compare text



Speech to text

Learn how to turn speech into text

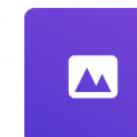


Image generation

Learn how to generate or edit images



Fine-tuning

Learn how to train a model for your use case

Exploring OpenAI Generative AI Models

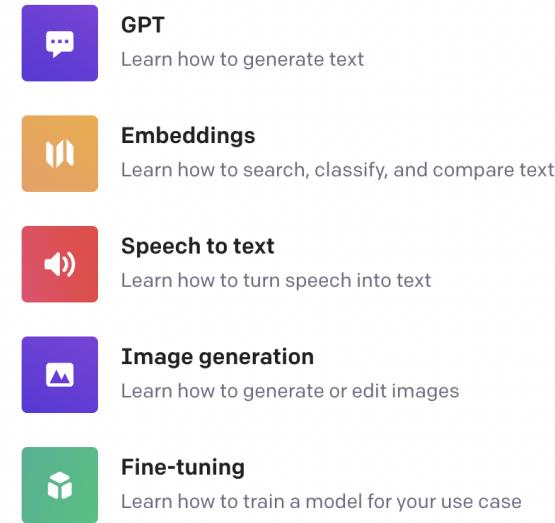
In 28
Minutes

Model	Notes
GPT-4	Generate natural language or code. Most advanced. gpt-4 (8,192 tokens), gpt-4-32k 32,768 tokens.
GPT-3.5	Generate natural language or code. gpt-3.5-turbo 4,097 tokens, gpt-3.5-turbo-16k 16,385 tokens.
DALL·E	Generate and edit images using natural language prompts
Whisper	Convert audio into text
Embeddings	Convert text into a numerical form
Moderation	Detect if text may be sensitive or unsafe (text-moderation-stable)
GPT-3 Legacy Models	Generate natural language or code. OLD. NOT RECOMMENDED (ada, babbage, curie, davinci).

Which text model to use?

In 28
Minutes

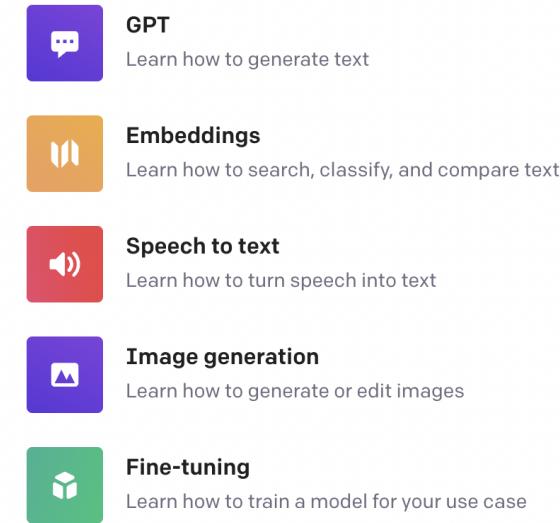
- Recommended: GPT-4 or GPT-3.5
- **GPT-4:**
 - Generally performs better at complex tasks
 - Generally less prone to hallucinations (making up information)
 - Has larger default context window (8,192 tokens)
- **GPT-3.5:**
 - Lower latency
 - Costs much less per token
- **Recommended:** Experiment for your use case



Understanding Pricing of OpenAI APIs

In 28
Minutes

- Updates: <https://openai.com/pricing>
- Open AI APIs
 - Prices are per 1K tokens
 - Prices vary from model to model
 - (MY EXPERIENCE) Prices vary with time
 - Start experimenting for free
- Both Open AI APIs and ChatGPT Plus are independent from each other
 - Access to ChatGPT Plus will NOT provide access to Open AI APIs



Exploring Open AI API - Basic Prompting

In 28
Minutes

```
import os
import openai

openai.api_key = os.getenv("OPENAI_API_KEY")

response = openai.ChatCompletion.create(
    model="gpt-*****",
    messages=[
        {
            "role": "user",
            "content": "Generate a list of Top 2 DevOps and Cloud Trends"
        }
    ],
    temperature=1,
    max_tokens=256,
    top_p=1,
    frequency_penalty=0,
    presence_penalty=0
)
```

Response Example

```
{  
    "id": "chatcmpl-*****",  
    "object": "chat.completion",  
    "created": 1693824659,  
    "model": "gpt-*****", #MODEL used for the chat completion  
    "choices": [  
        {  
            "index": 0,  
            "message": {  
                "role": "assistant",  
                "content": "AN EXAMPLE RESPONSE GOES HERE!" #RESPONSE  
            },  
            "finish_reason": "stop" #REASON the model stopped generating tokens  
                #(`stop` - model hit a natural stop point,  
                #`length` - max tokens reached,...)  
        }  
    ],  
    "usage": {  
        "prompt_tokens": 167, #No of tokens in the prompt  
        "completion_tokens": 47, #No of tokens in the generated completion  
        "total_tokens": 214 #Total No of tokens used in the request (prompt + completion).  
    }  
}
```

Exploring Open AI API - Prompts With Examples

In 28
Minutes

```
# Typical flow: System Message (Optional), alternating user and assistant messages
# System message: Set the behavior of the assistant

response = openai.ChatCompletion.create(
    model="gpt-*****",
    messages=[
        {
            "role": "system",
            "content": "CONTEXT GOES IN HERE"
        },
        {
            "role": "user",
            "content": "EXAMPLE-INPUT-1"
        },
        {
            "role": "assistant",
            "content": "EXAMPLE-OUTPUT-1"
        },
        {
            "role": "user",
            "content": "SAMPLE-REQUEST"
        }
    ], temperature=1, max_tokens=256, top_p=1 #others
)
```

Exploring Chat API

In 28
Minutes

- A number of enterprises provide chat bots for customers!
- You can easily create such chat bots using OpenAI API.
- **Usecases**
 - **Customer service:** Answer customer questions, troubleshoot issues, and provide information
 - **Education and training:** Assess the level of a student, answer questions, and give feedback.

The screenshot shows a chat application interface. On the left, there is a sidebar with the text: "SYSTEM" followed by three paragraphs of text describing the AI's role as an expert on the Solar System, its limitations regarding other topics, and its commitment to truthfulness. To the right, there are two separate conversations. The top conversation is between a "USER" who asks about Pluto and an "ASSISTANT" who provides a detailed response. The bottom conversation is between a "USER" who asks about AWS and an "ASSISTANT" who responds that it is outside the AI's expertise. At the bottom of the interface, there are buttons for "Submit" and "Add message", and a link to "Give us feedback".

SYSTEM

You are Mr Solar System, an expert on Astronomy. You know everything about Solar System. Respond in short sentences.

You do NOT know anything about topics other than Solar System

You are truthful and never lie. Never make up facts and if you are not 100% sure, reply with why you cannot answer in a truthful way.

USER Tell me about Pluto?

ASSISTANT Pluto is a dwarf planet in our Solar System. It used to be classified as the ninth planet, but in 2006 it was reclassified as a dwarf planet. Pluto is smaller than Earth's moon.

USER Yup. I was told Pluto was a planet when studying in school!

ASSISTANT Yes, that was the case for many years. However, the definition of a planet changed in 2006, and Pluto no longer met the criteria to be considered a planet.

USER Tell me about AWS

ASSISTANT I apologize, but I am an expert in Astronomy and the Solar System. I do not have knowledge about AWS (Amazon Web Services).

Add message

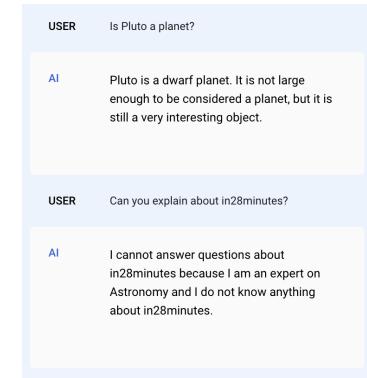
Submit

Give us feedback

Chat - Best Practices

In 28
Minutes

- **1: Assume a character**
 - You are Mr Solar System, an expert on Astronomy
- **2: Identify Topics to focus on**
 - You know everything about Solar System
- **3: Specify response style**
 - Respond in short sentences
 - Shape your response as if talking to a 10-years-old
- **4: Provide specific rules**
 - You do NOT know anything about topics other than Solar System
- **5: Add a rule to reduce hallucinations**
 - You are truthful and never lie. Never make up facts and if you are not 100% sure, reply with why you cannot answer in a truthful way.



Exploring Open AI API - Chat Example

In 28
Minutes

```
# Typical flow: System Message (Optional), alternating user and assistant messages
# System message: Set the behavior of the assistant
response = openai.ChatCompletion.create(
    model="gpt-*****",
    messages=[
        {
            "role": "system",
            "content": "PROVIDE CONTEXT HERE"
        },
        {
            "role": "user",
            "content": "SAMPLE QUESTION"
        },
        {
            "role": "assistant",
            "content": "SAMPLE ANSWER"
        },
        {
            "role": "user",
            "content": "YOUR NEXT QUESTION"
        }
    ],
    temperature=0.2,max_tokens=256 #others
)
```

Exploring Open AI API - Image & Audio Examples

In 28
Minutes

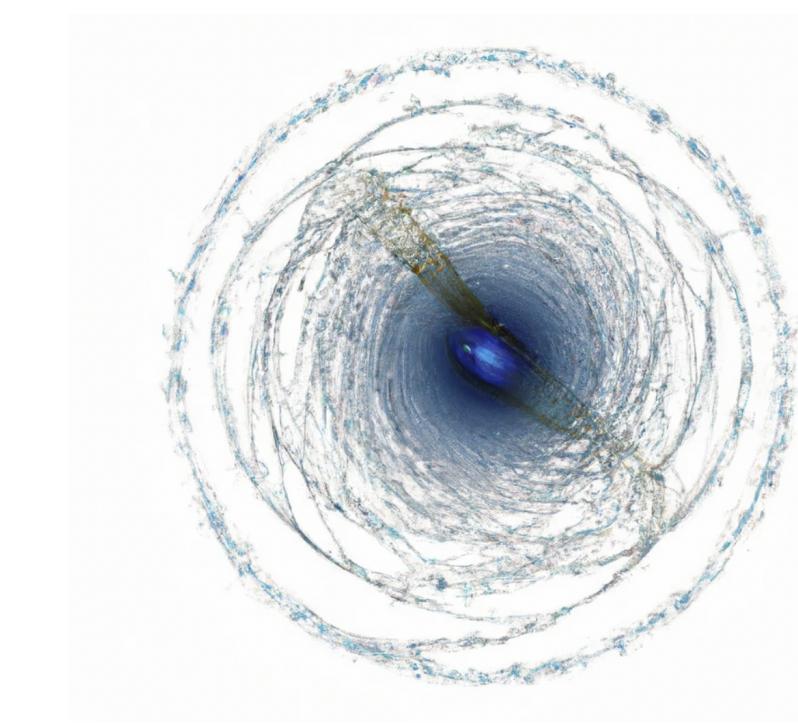
```
//Model - DALL·E
response = openai.Image.create(
    prompt="Your image prompt",
    n=1, #No of images to generate. Must be between 1 and 10.
    size="1024x1024" #256x256, 512x512, or 1024x1024
)
image_url = response[ 'data' ][0][ 'url' ]

audio_file= open("/path/to/file/audio.mp3", "rb")
transcript = openai.Audio.transcribe("whisper-1", audio_file)
transcript = openai.Audio.translate("whisper-1", audio_file)
```

Tuning Language Models

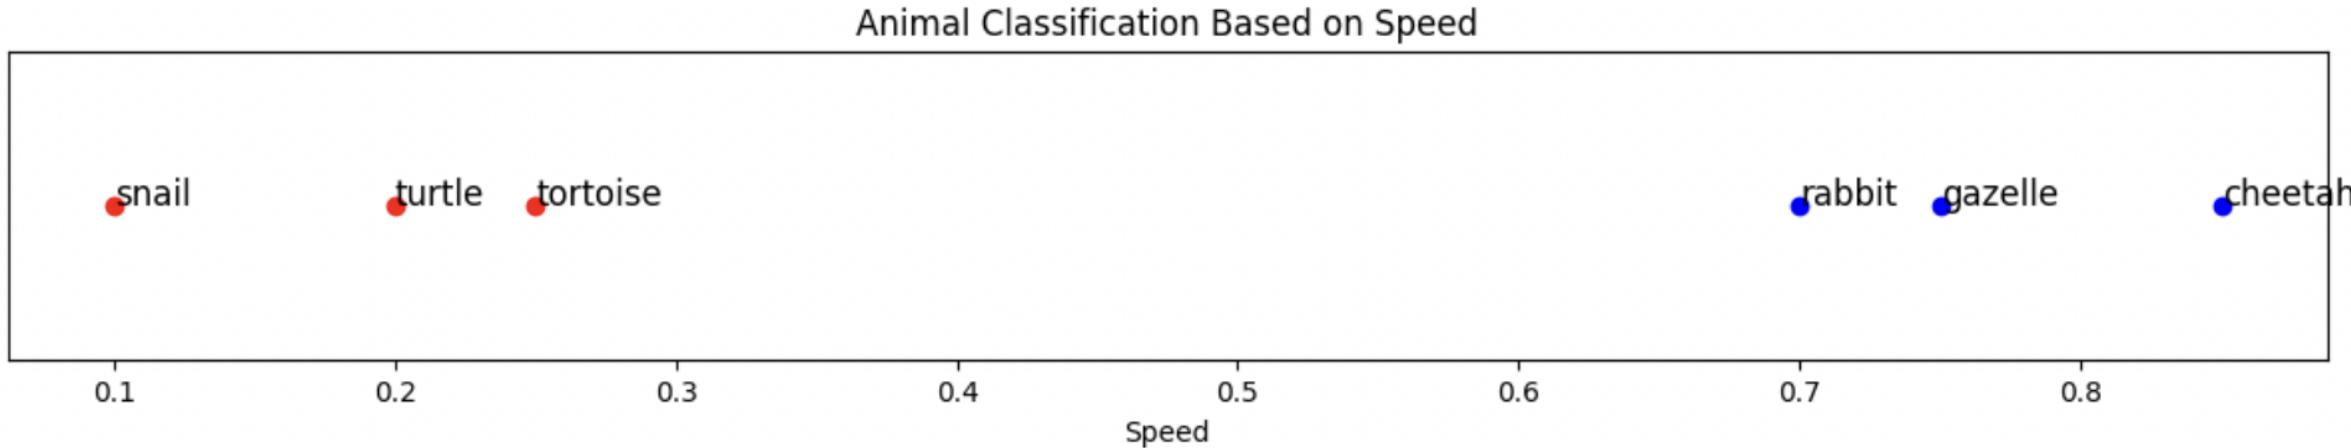
In 28
Minutes

- Sometimes text models might not give you sufficient accuracy
- In such scenarios, you can **Tune an LLM**
 - Example: Tune gpt-3.5-turbo LLM
- **Steps:**
 - 1: Create a training dataset file (JSONL format)
 - 2: Tune the model using the dataset
 - 3: Use the tuned model



What are Embeddings?

In 28
Minutes



- **Embeddings:** Vector representations of words in a high-dimensional space
 - Captures semantic relationships and contextual information
- **Example:** You can use multiple dimensions to represent animals:
 - Habitat: "aquatic," "terrestrial," or "arboreal."
 - Diet: "carnivore," "herbivore," or "omnivore."
 - Size: "small," "medium," or "large."
 - Movement: "flying," "running," "swimming," or "crawling."

Exploring Embeddings with an Example

In 28
Minutes

- On the right is an embedding of a single word
 - OpenAI Embeddings API provides 1536-dimensional vector embeddings
 - i.e. Each word is being looked at from 1536 different dimensions
- Widely used in natural language processing (NLP) tasks
 - **Text Similarity:** Measure semantic similarity between texts
 - **Recommendation Systems:** Recommend items based on user preferences
 - **Clustering:** Group similar texts
 - **Outlier Detection:** Find text that does not fit the group
 - **Example: Similarity Calculation**
 - Given two sentences
 - "The sun is shining brightly." and "Cats and dogs are popular pets."
 - Calculate similarity between the sentence embeddings.
 - Higher similarity indicates semantic closeness.

```
[0.00020168583432678133,  
 0.017162907868623734,  
 0.02314572036266327,  
 0.01056084968149662,  
 0.04190816730260849,  
 -0.02385203167796135,  
 -0.007645965088158846,  
 0.022990167140960693,  
 -0.026320133358240128,  
 0.02654663473367691,  
 -0.050877638161182404,  
 -0.006736526265740395,  
 0.009900923818349838,  
 0.00828093197196722,  
 -0.023270031437277794,  
 -0.052012279629707336,  
 -0.04786107689142227,  
 -0.020648762583732605,  
 -0.006686172913759947,  
 -0.0021143548656255007,  
 -0.05750234052538872,  
 -0.0331496000289917,  
 -0.03808722645044327,  
 -0.023742586374282837,  
 -0.006033886689692736,  
 -0.10131429135799408,  
 0.0332576185464859,  
 0.022916549816727638,  
 -0.0483529306948185,  
 -0.010218596085906029,  
 -0.0030683630611747503,  
 0.014686625450849533,
```

- **LangChain - Build flexible applications powered by LLMs**

- **1:** Abstractions for working with language models

- Easier to work with LLMs
 - Easily switch from one LLM to another
 - **LLM (Text) Model** - Represents pure text completion models (text in, text out)
 - VertexAI, GooglePalm, OpenAI, AzureOpenAI, LlamaCpp
 - **Chat Model** - Instead of “text in, text out”, exposes an interface where “chat messages” are the inputs and outputs
 - ChatVertexAI, ChatGooglePalm, ChatOpenAI, AzureChatOpenAI
 - **Embedding Model** - Interface for embedding models
 - OpenAIEmbeddings, VertexAIEmbeddings, GooglePalmEmbeddings, LlamaCppEmbeddings



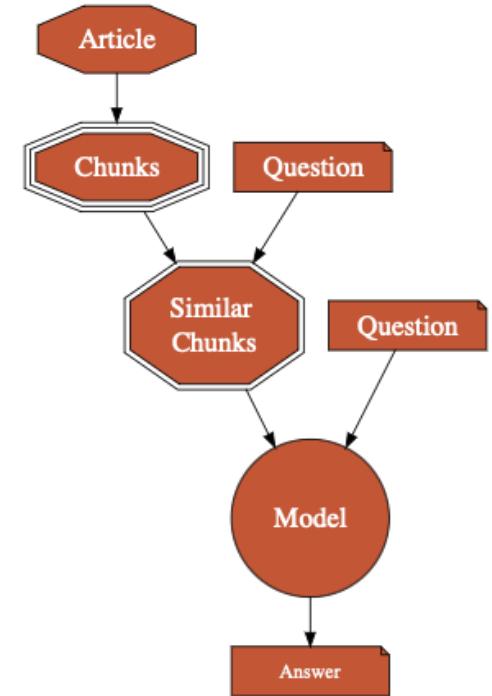
- **2:** Components for doing higher-level tasks

- Define a sequence of steps as a chain
 - Answer questions looking at information from different sources
 - Summarize long pieces of text

LangChain - Long Article - Q & A

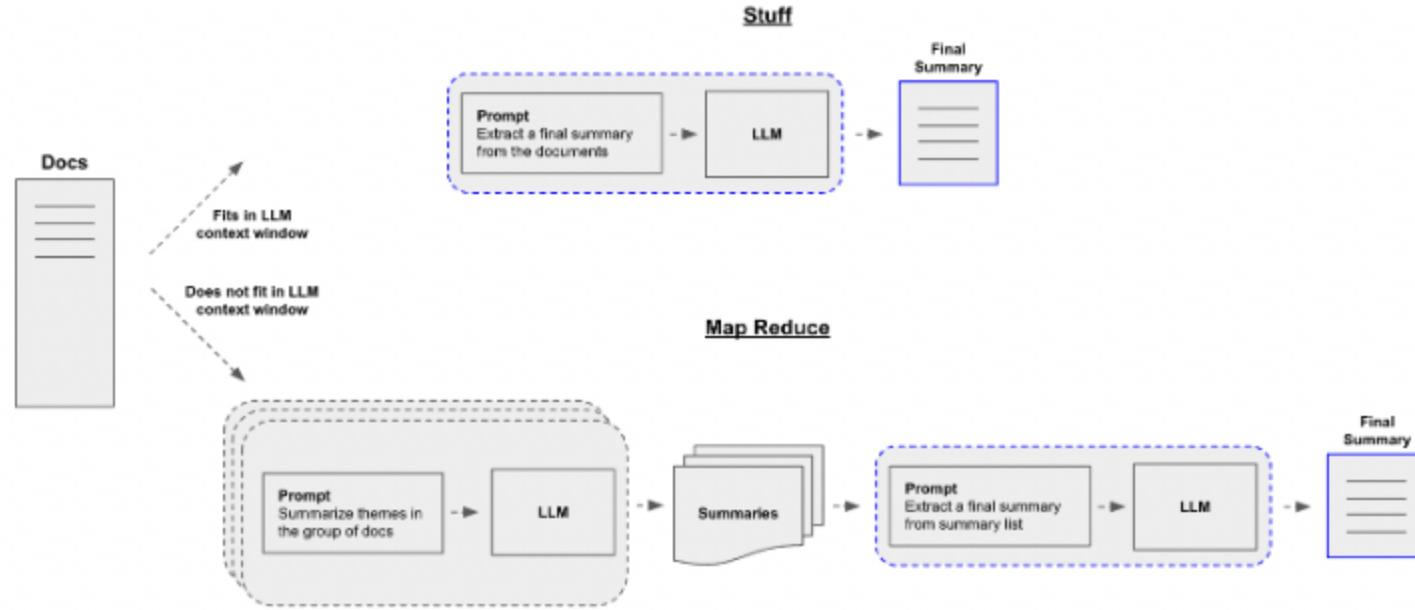
In 28
Minutes

- I want to answer questions based on an Article
- **Approach 1:** Use Article + Question as Prompt
 - Simplest Approach
 - **PROBLEM:** Does NOT work with Long Articles
- **Approach 2:** Split article into multiple parts > Find parts where question is answered
 - Complex Approach
 - 1: Split Article into multiple chunks
 - RecursiveCharacterTextSplitter
 - 2: Do a Similarity Search of question with each chunk
 - FAISS + Embeddings
 - 3: Use Similar Chunk(s) + Question as Prompt



LangChain - Summarization

In 28
Minutes

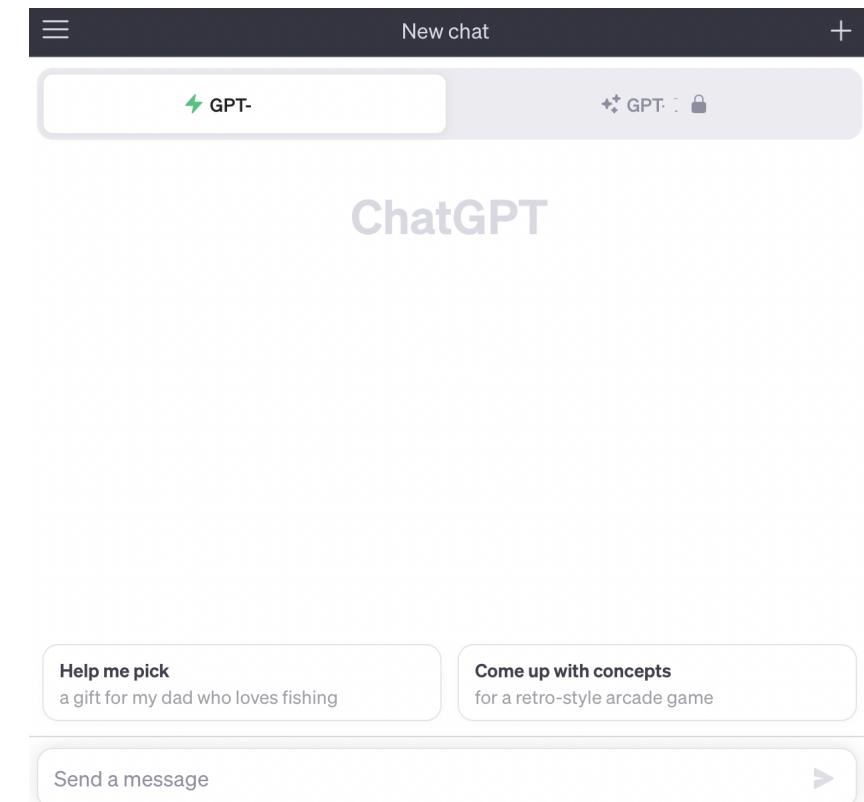


- Summarize content from multiple sources
 - **Stuff:** "stuff" all documents into a single prompt
 - Simplest Approach HOWEVER might not be feasible for long documents
 - **Map-reduce:** Summarize each document first
 - AND then "reduce" the summaries into a final summary

Exploring ChatGPT Plans

In 28
Minutes

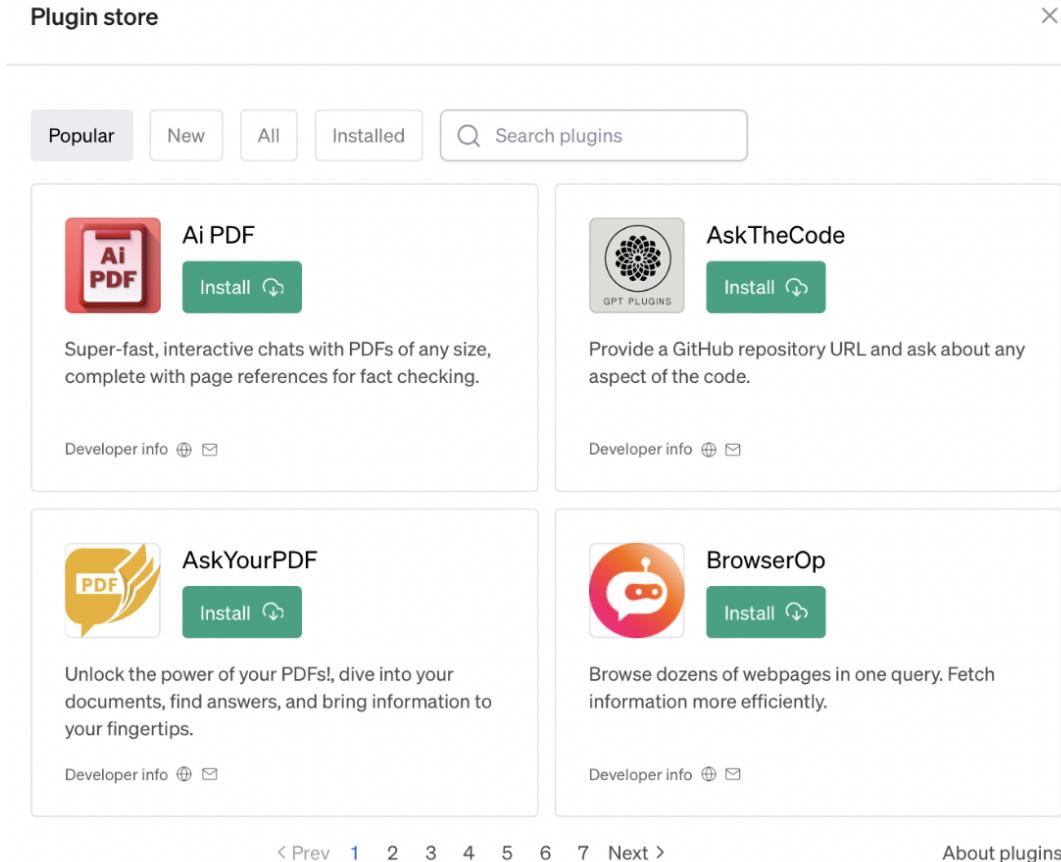
- REMEMBER: These plans will change with time
 - Latest: <https://help.openai.com/en/collections/3742473-chatgpt>
 - ChatGPT Free: GPT-3.5 access
 - ChatGPT Plus: Paid Subscription Plan
 - GPT-4 access (within a few limits)
 - Faster response speed
 - Priority Access to New Features
 - Plugins
 - Advanced Data Analysis
 - REMEMBER: OpenAI API access is NOT included
 - ChatGPT Enterprise: Enterprise-grade security and privacy
 - Unlimited higher-speed GPT-4 access
 - Longer context windows (32k token context window)
 - Shareable chat templates (across an enterprise)



Getting Started with ChatGPT Plugins

In 28
Minutes

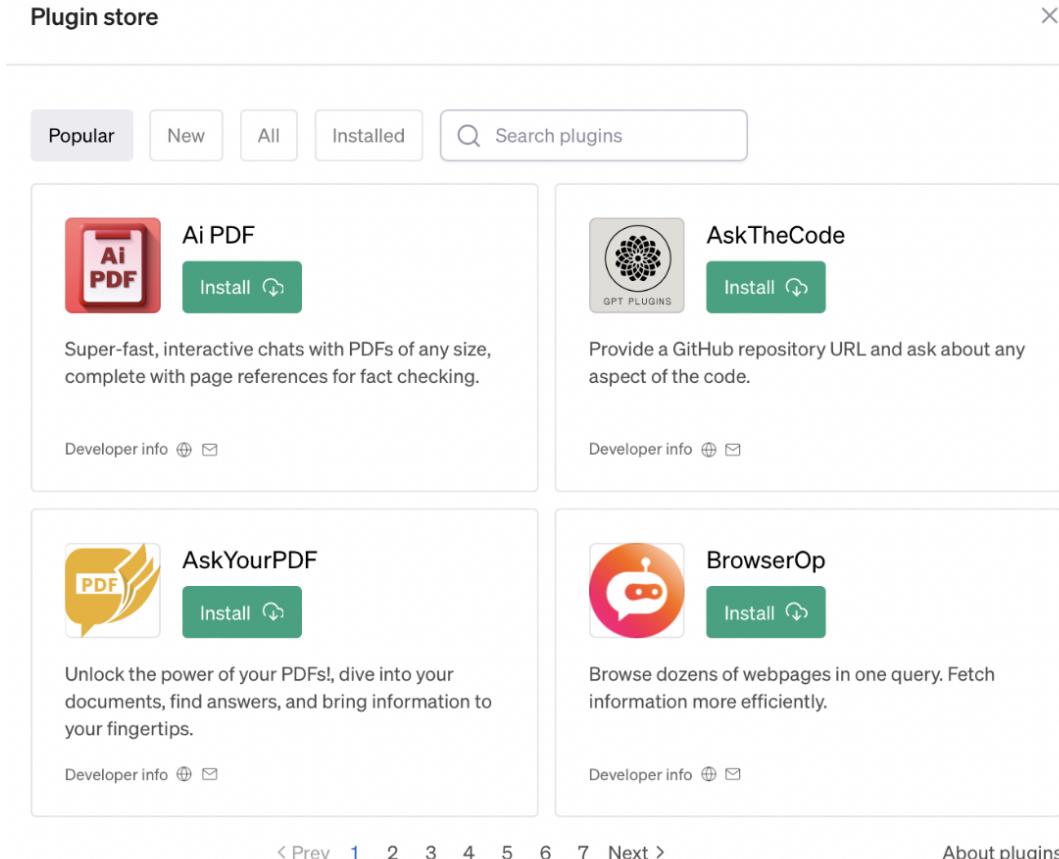
- **ChatGPT Plugins:** Add amazing new features to ChatGPT
- Powered by **third party** applications
 - NOT controlled by OpenAI
 - Only enable those you trust
- **REMEMBER:** ChatGPT may send your conversation and location (country or state) you're in to the plugin
- ChatGPT determines when to use **plugins during a conversation**
 - Depending on the plugins you've enabled
- Paid Feature (Needs to be Enabled!)



Exploring ChatGPT Plugins

In 28
Minutes

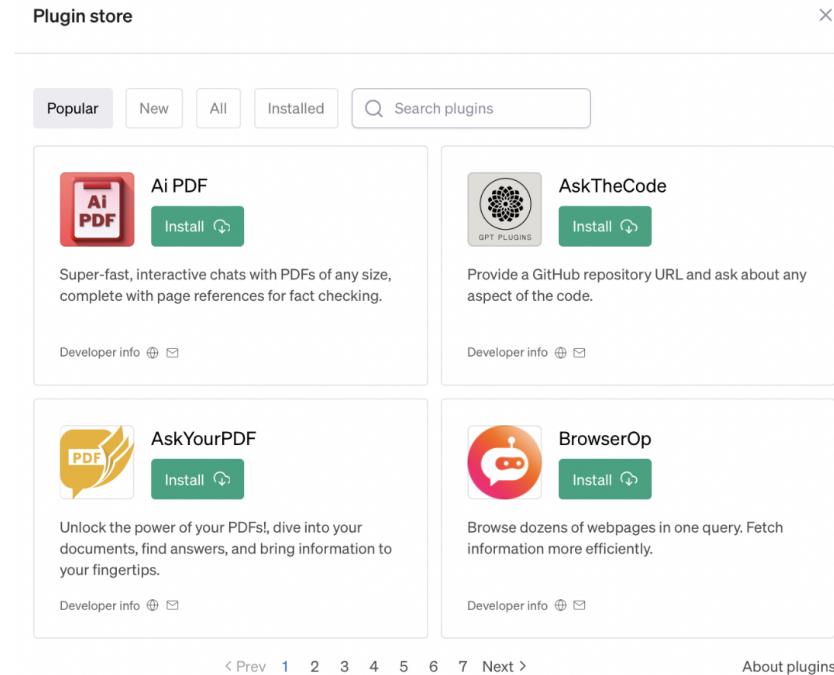
- **Wikipedia** - Tell me about 2028 Summer Olympics
- **Kayak** - I want to fly from New Delhi to New York tomorrow. Can you suggest the best flight?
- **Link Reader** - Can you summarise content from this page - <https://www.udemy.com/user/in28minutes>
- **Prompt Perfect** - perfect - tell me about in28minutes



Exploring ChatGPT Plugins - 2

In 28
Minutes

- **daigr.am** - Create a line chart showing the average global temperature for each decade from 1900 to 2020.
- **Diagrams:Show Me** - Create a diagram explaining how a HTTP request is handled
 - show ideas
 - Add More Details and Include Error Handling
- **Make a Sheet** - Create a CSV file with details of Top 10 countries by population - Include capital city, area and population as columns.
- **ScholarAI** - Can you tell me about some latest research on ChatGPT models?

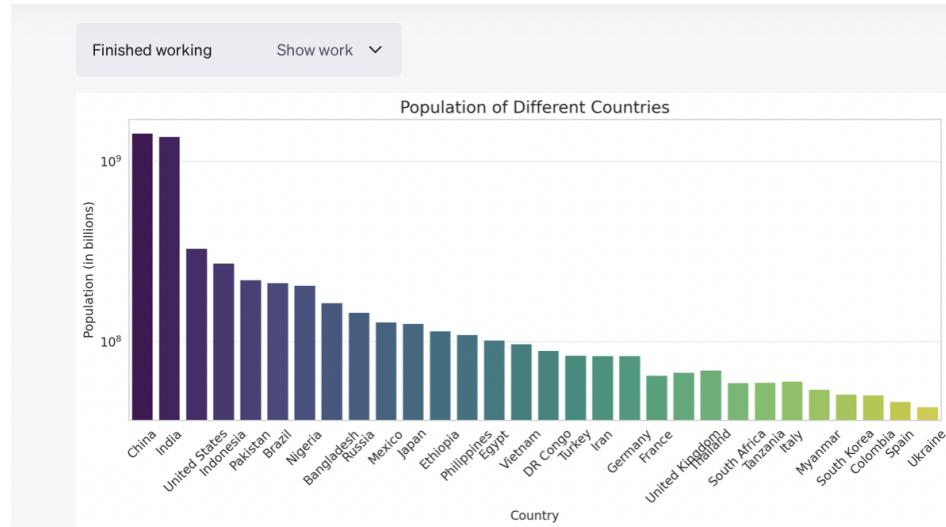


Exploring ChatGPT Advanced Data Analysis

In 28
Minutes

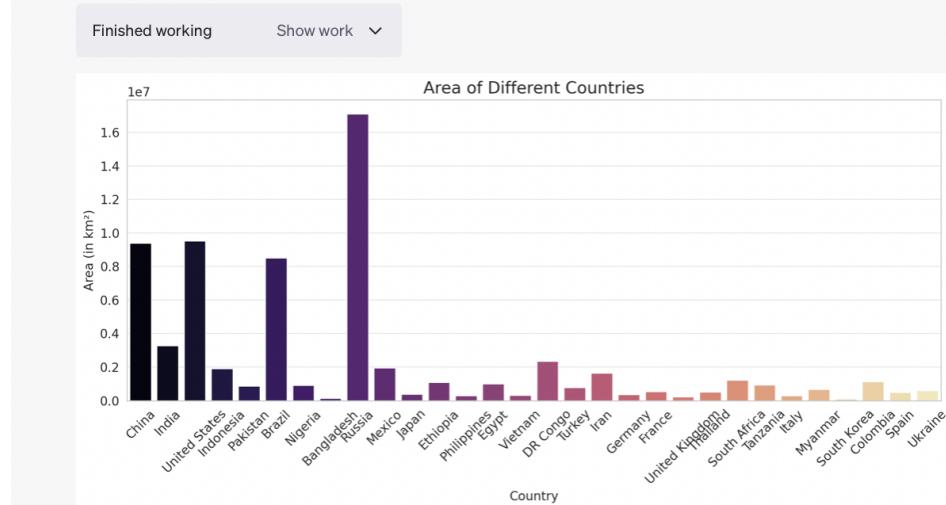
- **"Analyze information in seconds"**
(Verbatim - OpenAI announcement)

- Earlier called "Code Interpreter"
- Uses Python code
- Paid Feature (Needs to be Enabled!)
- ***Let's Play:***
 - Can you describe the data?
 - Create few visualizations to understand the data
 - Add a row to CSV with population density
 - Show me 5 example rows please
 - Calculate Density to 1 decimal point accuracy
 - Find top 5 countries by population density
 - Can you show this image?
 - Can you convert it to grayscale?
 - Can you reduce the height and width by 50% each?
 - What image formats have smaller size?



Here's a bar chart comparing the populations of different countries. The y-axis is on a logarithmic scale to better capture the vast differences in population numbers.

Next, let's create a bar chart comparing the areas (in km²) of different countries.

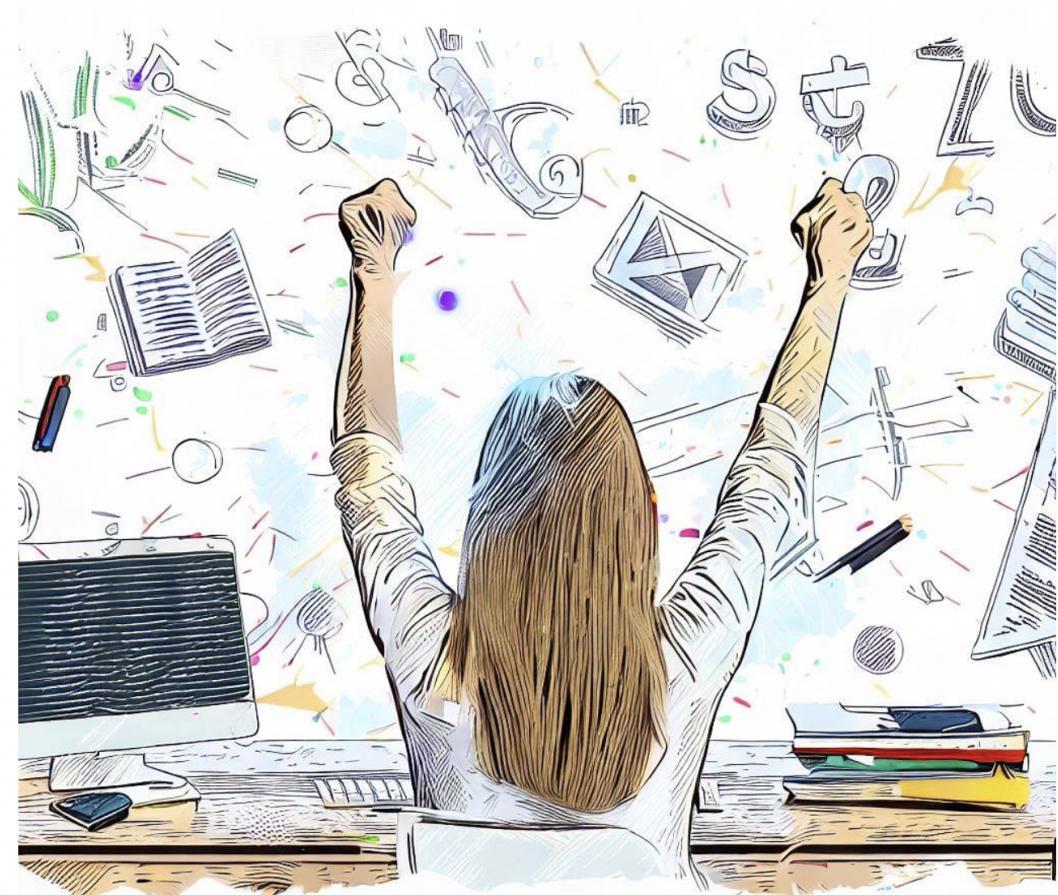


Get Ready

Let's clap for you!

In 28 Minutes

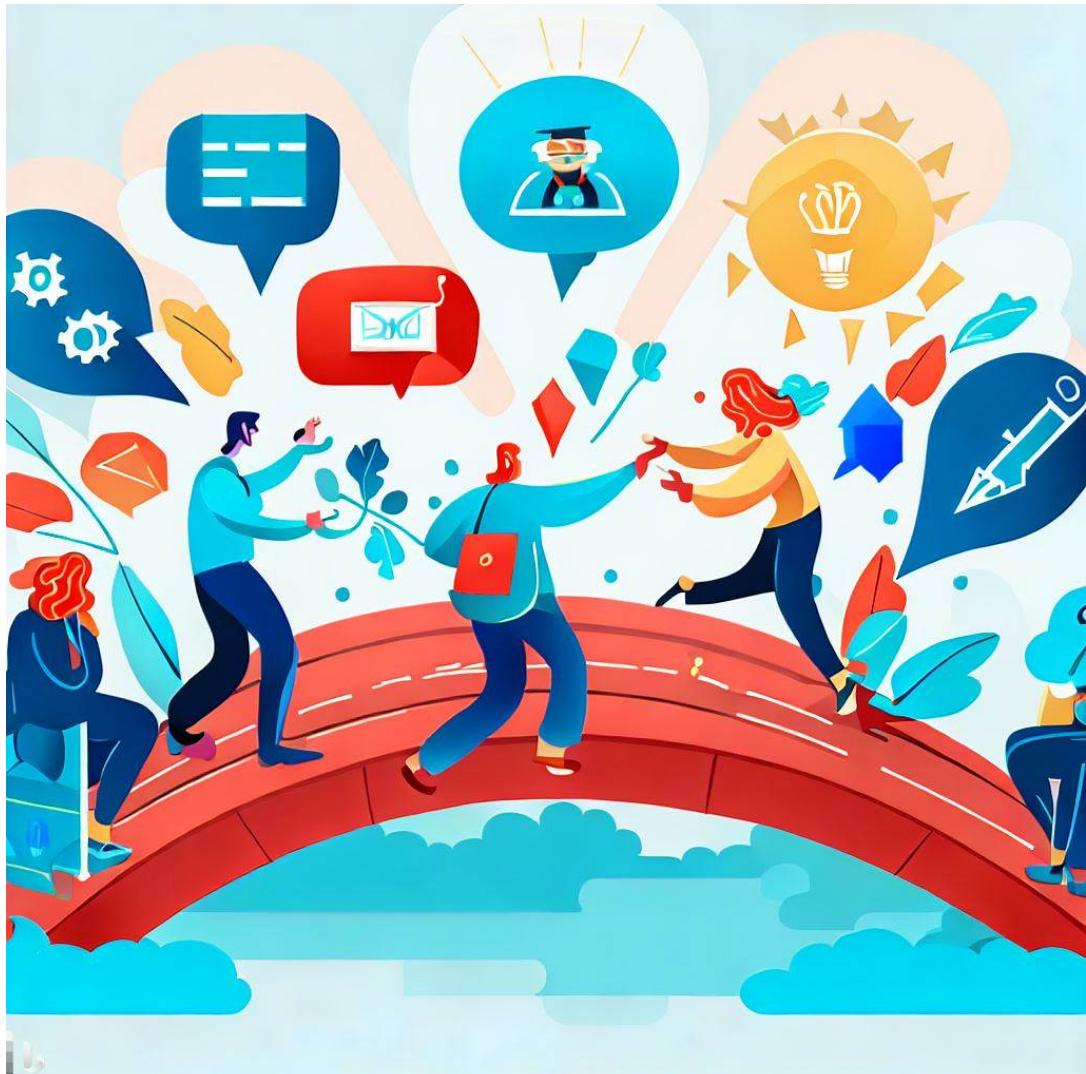
- Congratulations
 - You have put your best foot forward to start your Generative AI Journey!
 - Good Luck!
 - Keep Learning Every Day!



Do Not Forget

In 28
Minutes

- Recommend the course to your friends!
 - Do not forget to review!
- Your Success = My Success
 - Share your success story with me on LinkedIn (Ranga Karanam)
 - Share your success story and lessons learnt in Q&A with other learners!



What next?

In 28
Minutes

- Learn Other Cloud Platforms:
 - Gartner predicts a multi cloud world soon
 - Get certified on AWS, Azure and Google Cloud
- Learn DevOps (Containers and Container Orchestration)
- Learn Full Stack Development



